

Hash Tables

Complexity using recurrence

- Consider an algorithm which is recursive and assume it takes $T(n)$ as the complexity
- For binary search, we divide into two equal parts and we work on one of the
- Assume the it takes $f(n)=O(1)$ to combine the results
- Then, $T(n)=T(n/2)+f(n)$
- If $f(n)=1$, then
- $T(n) = 1 + T(n/2) = 1 + (1 + T(n/4)) = 2 + T(n/4)$
 $= 2 + (1 + T(n/8)) = 3 + T(n/8) = \dots$
 $k + T(n/2^k) = \log n + T(n/2^{\log n}) = \log n + T(1) =$
 $\log n + 1 = O(\log n).$

Master theorem for complexity


- **General formula that works if recurrence has the form**
- **$T(n) = aT(n/b) + f(n)$**
 - **a is number of subproblems**
 - **n/b is size of each subproblem**
 - **$f(n)$ is cost of non-recursive part**
- a) **If $f(n) = O(n^{\log_b(a)-\epsilon})$, then $T(n) = \Theta(n^{\log_b(a)})$.**
- b) **If $f(n) = \Theta(n^{\log_b(a)})$, then $T(n) = \Theta(n^{\log_b(a)} \log(n))$.**
- c) **If $f(n) = \Omega(n^{\log_b(a)+\epsilon})$ and then $T(n) = \Theta(f(n))$**

Hash Tables

- All search structures so far
 - Relied on a comparison operation
 - Performance $O(n)$ or $O(\log n)$
- Assume I have a function
 - $f(key) \rightarrow integer$
ie one that maps a key to an integer
- What performance might I expect now?

Hash Tables - Structure

- **Simplest case:**
 - Assume items have integer keys in the range $1 \dots m$
 - Use the value of the key itself to select a slot in a **direct access table** in which to store the item
 - To search for an item with key, k , just look in slot k
 - If there's an item there, you've found it
 - If the tag is 0, it's missing.
 - **Constant time, $O(1)$**




Hash Tables - Constraints

- **Constraints**
 - Keys must be unique
 - Keys must lie in a small range
 - For storage efficiency, keys must be **dense** in the range
 - If they're **sparse** (lots of gaps between values), a lot of space is used to obtain speed
 - **Space for speed trade-off**

Hash Tables - Relaxing the constraints


- **Keys must be unique**
 - Construct a linked list of duplicates “attached” to each slot
 - If a search can be satisfied by *any* item with key, k , performance is still $O(1)$
but
 - If the item has some other distinguishing feature which must be matched, we get $O(n^{max})$

where n^{max} is the largest number of duplicates - or length of the longest chain



Hash Tables - Relaxing the constraints

- Keys are integers
 - Need a **hash function**
 $h(\text{key}) \rightarrow \text{integer}$
ie one that maps a key to an integer
 - Applying this function to the key produces an address
 - If h maps each key to a **unique integer** in the range $0 \dots m-1$ then search is $O(1)$



Hash Tables - Hash functions

- Form of the hash function

- Example - using an n -character key

```
int hash( char *s, int n ) {  
    int sum = 0;  
    while( n-- ) sum = sum + *s++;  
    return sum % 256;  
}
```

returns a value in 0 .. 255

- xor function is also commonly used

```
sum = sum ^ *s++;
```

- But **any** function that generates integers in $0..m-1$ for some suitable (*not too large*) m will do
 - As long as the hash function itself is $O(1)$!

Hash Tables - Collisions

- Hash function

- With this hash function

```
int hash( char *s, int n ) {  
    int sum = 0;  
    while( n-- ) sum = sum + *s++;  
    return sum % 256;  
}
```

- hash("AB", 2) and
hash("BA", 2)
return the same value!
 - This is called a collision
 - A variety of techniques are used for resolving collisions

Hash Tables - Collision handling

- **Collisions**

- Occur when the hash function maps two **different keys** to the **same address**
- The table must be able to recognise and resolve this
- **Recognise**
 - Store the actual key with the item in the hash table
 - Compute the address
 - $k = h(\text{key})$
 - Check for a hit
 - *if (table[k].key == key) then **hit***
*else **try next entry***

- **Resolution**

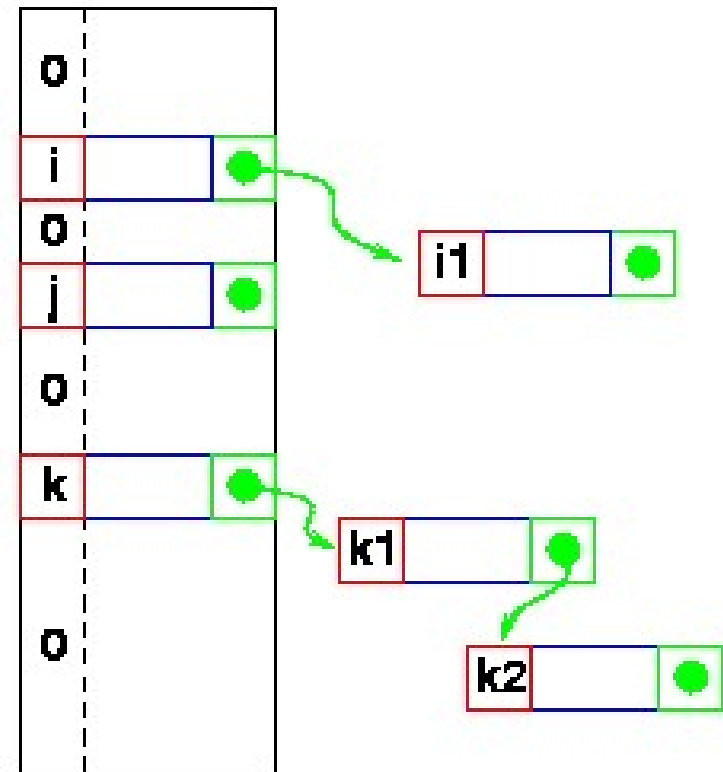
- **Variety of techniques**



**We'll look at various
“try next entry” schemes**

Hash Tables - Linked lists

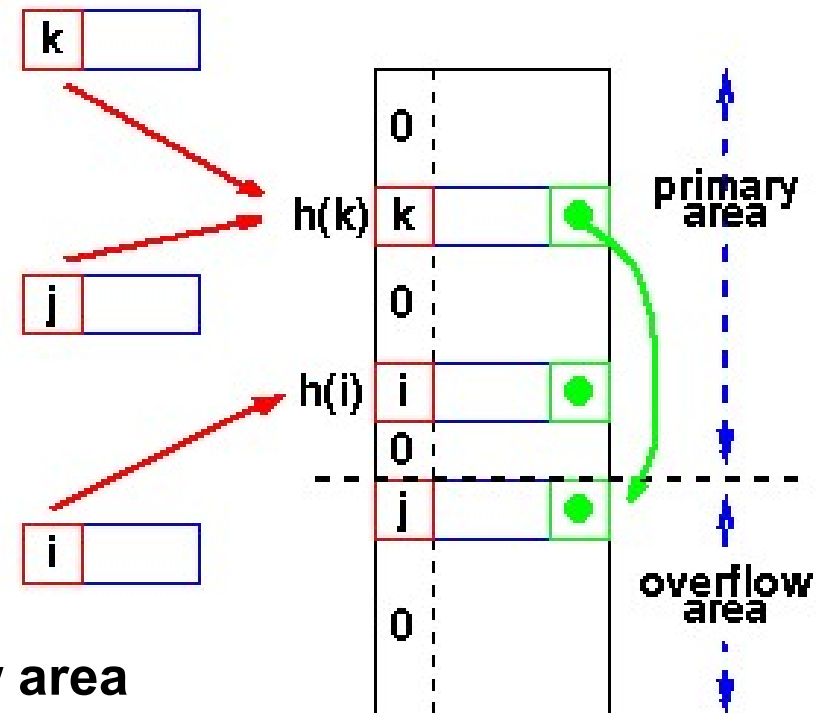
- **Collisions - Resolution**
 - **Linked list attached to each primary table slot**
 - $h(i) == h(i1)$
 - $h(k) == h(k1) == h(k2)$
 - **Searching for $i1$**
 - Calculate $h(i1)$
 - Item in table, i , doesn't match
 - Follow linked list to $i1$
 - **If NULL found, key isn't in table**



Hash Tables - Overflow area

□ Overflow area


- Linked list constructed in special area of table called **overflow area**
- $h(k) == h(j)$
- **k** stored first
- Adding **j**
 - Calculate $h(j)$
 - Find **k**
 - Get first slot in overflow area
 - Put **j** in it
 - **k**'s pointer points to this slot
- Searching - same as linked list



Hash Tables - Re-hashing

□ Use a second hash function

- Many variations
- General term: **re-hashing**
- $h(k) == h(j)$
- **k** stored first
- Adding **j**
 - Calculate $h(j)$
 - Find **k**
 - Repeat for i until we find an empty slot
 - Calculate $h'(j, i)$
 - Put **j** in it
- Searching - Use $h(x)$, then $h'(x)$




**$h'(x)$ -
second hash function**

Hash Tables - Re-hash functions

□ The re-hash function

- Many variations
- **Linear probing**
 - $h'(x)$ is $+1$
 - Go to the next slot until you find one empty
- Can lead to bad **clustering**
- Re-hash keys fill in gaps between other keys and exacerbate the collision problem



Hash Tables - Re-hash functions

□ The re-hash function

- Many variations
- Quadratic probing
 - $h'(x)$ is $c i^2$ on the i^{th} probe
 - Avoids primary clustering
 - Secondary clustering occurs
 - All keys which collide on $h(x)$ follow the same sequence
 - First
 - $a = h(j) = h(k)$
 - Then $a + c, a + 4c, a + 9c, \dots$
 - Secondary clustering generally less of a problem

Hash Tables - Collision Resolution Summary

- **Chaining**
 - + Unlimited number of elements
 - + Unlimited number of collisions
 - Overhead of multiple linked lists
- **Re-hashing**
 - + Fast re-hashing
 - + Fast access through use of main table space
 - Maximum number of elements must be known
 - Multiple collisions become probable
- **Overflow area**
 - + Fast access
 - + Collisions don't use primary table space
 - Two parameters (size/overflow size) which govern performance need to be estimated

Hash Tables - Summary so far ...

- Potential $O(1)$ search time
 - If a suitable function $h(key) \rightarrow integer$ can be found
- Space for speed trade-off
 - “Full” hash tables don’t work (more later!)
- Collisions
 - Inevitable
 - Hash function reduces amount of information in key
 - Various resolution strategies
 - Linked lists
 - Overflow areas
 - Re-hash functions
 - Linear probing h' is $+1$
 - Quadratic probing h' is $+ci^2$
 - Any other hash function!
 - or even sequence of functions!

Hash Tables - Choosing the Hash Function

- **“Almost any function will do”**
 - **But some functions are definitely better than others!**
- **Key criterion**
 - **Minimum number of collisions**
 - **Keeps chains short**
 - **Maintains $O(1)$ average**

Hash Tables - Choosing the Hash Function

- Uniform hashing
 - Ideal hash function
 - $P(k)$ = probability that a key, k , occurs
 - If there are m slots in our hash table,
 - a **uniform hashing function**, $h(k)$, would ensure:

$$\sum_{k \mid h(k) = 0} P(k) = \sum_{k \mid h(k) = 1} P(k) = \dots \sum_{k \mid h(k) = m-1} P(k) = \frac{1}{m}$$

Read as sum over all k such that $h(k) = 0$

- *or, in plain English,*
- the number of keys that map to each slot is equal

Hash Tables - A Uniform Hash Function

- *If the keys are integers randomly distributed in $[0, r)$, then*

$$h(k) = \left\lfloor \frac{mk}{r} \right\rfloor$$



Read as $0 \leq k < r$

is a **uniform hash function**

- Most hashing functions can be made to map the keys to $[0, r)$ for some r
 - *eg adding the ASCII codes for characters mod 255 will give values in $[0, 256)$ or $[0, 255]$*
 - **Replace + by xor**
 - **same range without the mod operation**

Hash Tables - Reducing the range to [0, m)

- **We've mapped the keys to a range of integers**
 $0 \leq k < r$
- **Now we must reduce this range to [0, m)**
where m is a reasonable size for the hash table
- **Strategies**
 - **Division - use a mod function**
 - **Multiplication**
 - **Universal hashing**

Hash Tables - Reducing the range to $[0, m)$

□ Division

- Use a mod function

$$h(k) = k \bmod m$$

- Choice of m ?

- Powers of 2 are generally not good!

$$h(k) = k \bmod 2^n$$

selects **last n bits of k**

$k \bmod 2^8$ selects these bits

0110010111000011010

- All combinations are not generally equally likely
- **Prime numbers close to 2^n seem to be good choices**
eg want ~4000 entry table, choose $m = 4093$

Hash Tables - Reducing the range to $[0, m)$

□ Multiplication method

- Multiply the key by constant, A , $0 < A < 1$
- Extract the fractional part of the product

$$(kA - \lfloor kA \rfloor)$$

- Multiply this by m

$$h(k) = \lfloor m * (kA - \lfloor kA \rfloor) \rfloor$$

- Now m is not critical and a power of 2 can be chosen
- So this procedure is fast on a typical digital computer
 - Set $m = 2^p$
 - Multiply k (w bits) by $\lfloor A \cdot 2^w \rfloor$ ζ $2w$ bit product
 - Extract p most significant bits of lower half
 - $A = \frac{1}{2}(\sqrt{5} - 1)$ seems to be a good choice (*see Knuth*)

Hash Tables - Reducing the range to $(0, m]$

□ Universal Hashing

- A determined “adversary” can always find a set of data that will defeat any hash function
 - Hash all keys to same slot $\hookrightarrow O(n)$ search
- Select the hash function randomly (at run time) from a set of hash functions
- -----
- Functions are selected *at run time*
 - Each run can give different results
 - *Even with the same data*
 - Good average performance obtainable

Collision Frequency

- **Birthdays *or* the von Mises paradox**
 - **There are 365 days in a normal year**
 - Birthdays on the same day unlikely?
 - **How many people do I need before “it’s an even bet”**
(ie the probability is $> 50\%$)
that two have the same birthday?
 - **View**
 - the days of the year as the slots in a hash table
 - the “birthday function” as mapping people to slots
 - **Answering von Mises’ question answers the question about the probability of collisions in a hash table**



Distinct Birthdays

- Let $Q(n)$ = probability that n people have distinct birthdays
- $Q(1) = 1$
- With two people, the 2nd has only 364 “free” birthdays


$$Q(2) = Q(1) * \frac{364}{365}$$

- The 3rd has only 363, and so on:

$$Q(n) = Q(1) * \frac{364}{365} * \frac{363}{365} * \dots * \frac{365-n+1}{365}$$

Coincident Birthdays

- Probability of having two identical birthdays
- $P(n) = 1 - Q(n)$
- $P(23) = 0.507$
- With 23 entries, table is only $23/365 = 6.3\%$ full!



Hash Tables - Load factor

- Collisions are very probable!

- Table load factor

$$\alpha = \frac{n}{m}$$

n = number of items

m = number of slots

must be kept low

- Detailed analyses of the average chain length (or number of comparisons/search) are available

- **Separate chaining**

- linked lists attached to each slot

gives best performance

- but uses more space!

- Resizing can be done to increase the size of the hash table

Hash Tables - General Design

□ Choose the table size

- Large tables reduce the probability of collisions!
- Table size, m
- n items
- Collision probability $\alpha = n / m$

□ Choose a table organisation

- Does the collection keep growing?
 - Linked lists (..... but consider a tree!)
- Size relatively static?
 - Overflow area *or*
 - Re-hash

□ Choose a hash function



Hash Tables - General Design

- **Choose a hash function**
 - A simple (and fast) one may well be fine ...
 - Read your text for some ideas!
- **Check the hash function against your data**
 - **Fixed data**
 - Try various h, m until the maximum collision chain is acceptable
 - Known performance
 - **Changing data**
 - Choose some representative data
 - Try various h, m until collision chain is OK
 - Usually predictable performance

Hash Tables - Review

- ***If you can meet the constraints***
- + **Hash Tables will generally give good performance**
- + **$O(1)$ search**