



Introduction to Relational Model

Union Operation

- The **union** operation allows us to combine two relations
- Notation: $r \cup s$
- For $r \cup s$ to be valid:
 - r, s must have the **same arity** (same number of attributes)
 - The attribute domains must be **compatible** (example: 2nd column of r deals with the same type of values as does the 2nd column of s)
- Example: To find all courses taught in the Fall 2017 semester, or in the Spring 2018 semester, or in both

$$\Pi_{course_id} (\sigma_{semester="Fall" \wedge year=2017}(section)) \cup \Pi_{course_id} (\sigma_{semester="Spring" \wedge year=2018}(section))$$

- Result:

| <i>course_id</i> |
|------------------|
| CS-101 |
| CS-315 |
| CS-319 |
| CS-347 |
| FIN-201 |
| HIS-351 |
| MU-199 |
| PHY-101 |

Set-Intersection Operation

- The **set-intersection** operation allows us to find tuples that are in both the input relations
- Notation: $r \cap s$
- Remember: $r \cap s = r - (r - s)$
- Assume:
 - r, s have the *same arity*
 - Attributes of r and s are compatible
- Example: Find the set of all courses taught in both the Fall 2017 and the Spring 2018 semesters

$\Pi_{course_id} (\sigma_{semester="Fall" \wedge year=2017}(section)) \cap \Pi_{course_id} (\sigma_{semester="Spring" \wedge year=2018}(section))$

- Result:

| <i>course_id</i> |
|------------------|
| CS-101 |

Set-Intersection Operation

- Relation r, s

| A | B |
|----------|-----|
| α | 1 |
| α | 2 |
| β | 1 |

| A | B |
|----------|-----|
| α | 2 |
| β | 3 |

- $r \cap s$

| A | B |
|----------|-----|
| α | 2 |

Set Difference Operation

- The **set-difference** operation allows us to find tuples that are in one relation but are not in another
- Notation $r - s$
- Set differences must be taken between **compatible** relations
 - r and s must have the **same** arity
 - Attribute domains of r and s must be compatible
- Example: Find all courses taught in the Fall 2017 semester, but not in the Spring 2018 semester

$$\Pi_{course_id} (\sigma_{semester="Fall" \wedge year=2017}(section)) - \Pi_{course_id} (\sigma_{semester="Spring" \wedge year=2018}(section))$$

- Result:

| <i>course_id</i> |
|------------------|
| CS-347 |
| PHY-101 |

Set Difference Operation

- Relation r, s

| A | B |
|----------|-----|
| α | 1 |
| α | 2 |
| β | 1 |

| A | B |
|----------|-----|
| α | 2 |
| β | 3 |

- $r - s$

| A | B |
|----------|-----|
| α | 1 |
| β | 1 |

The Assignment Operation

- It is convenient at times to write a relational-algebra expression by assigning parts of it to temporary relation variables
- The **assignment** operation is denoted by \leftarrow and works like assignment in a programming language
- Example: Find all ***instructor*** in the “Physics” and “Music” department

$$Physics \leftarrow \sigma_{dept_name="Physics"}(instructor)$$
$$Music \leftarrow \sigma_{dept_name="Music"}(instructor)$$
$$Physics \cup Music$$

- With the assignment operation, a query can be written as a sequential program consisting of a series of assignments followed by an expression whose value is displayed as the result of the query

The Rename Operation

- The results of relational-algebra expressions do not have a name that we can use to refer to them
- The **rename** operator ρ is provided for that purpose
- The expression returns the result of expression E under the name x

$$\rho_x(E)$$

- Another form of the rename operation:

$$\rho_{x(A1, A2, \dots, A_n)}(E)$$

- Example:

$$\rho_{Teacher(EmpID, EmpName)} \pi_{ID, Name}(\sigma_{Condition}(instructor))$$

Aggregate Operators

- **Aggregation function** takes a collection of values and returns a single value as a result

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values

- **Aggregate operation** in relational algebra

$$G_1, G_2, \dots, G_n \quad g \quad F_1(A_1), F_2(A_2), \dots, F_n(A_n) (E)$$

- E is any relational-algebra expression
- G_1, G_2, \dots, G_n is a list of attributes on which to group (can be empty)
- Each F_i is an aggregate function
- Each A_i is an attribute name

Aggregate Operators

- Relation r

| <i>A</i> | <i>B</i> | <i>C</i> |
|-----------------|-----------------|-----------------|
| α | α | 7 |
| α | β | 7 |
| β | β | 3 |
| β | β | 10 |

- $g_{sum(C)}(r)$

| <i>Sum-C</i> |
|---------------------|
| 27 |

Aggregate Operators

- Relation **account** grouped by **branch-name**

- $\text{branch-name } g \text{ sum(balance) (account)}$

| branch-name | Account-number | balance |
|--------------------|-----------------------|----------------|
| Perryridge | A - 102 | 400 |
| Perryridge | A - 201 | 900 |
| Brighton | A - 217 | 750 |
| Brighton | A - 215 | 750 |
| Redwood | A - 222 | 700 |

| branch-name | balance |
|--------------------|----------------|
| Perryridge | 1300 |
| Brighton | 1500 |
| Redwood | 700 |

- Result of aggregation does not have a name
 - Can use rename operation to give it a name
 - For convenience, we permit renaming as part of aggregate operation

$\text{branch-name } g \text{ sum(balance) as sum-balance (account)}$

Null Values

- It is possible for tuples to have a *null* value, denoted by null, for some of their attributes
- *null* signifies an unknown value or that a value does not exist
- The result of any arithmetic expression involving *null* is *null*
- Aggregate functions simply ignore *null* values
 - Is an arbitrary decision
 - Could have returned *null* as result instead
 - We follow the semantics of SQL in its handling of *null* values
- For duplicate elimination and grouping, *null* is treated like any other value, and two *nulls* are assumed to be the same
 - Alternative: Assume each *null* is different from each other
 - Both are arbitrary decisions, so we simply follow SQL

Null Values

- Comparisons with null values return the special truth value *unknown*
 - If *false* was used instead of *unknown*, then $\text{not } (A < 5)$ would not be equivalent to $A \geq 5$
- Three-valued logic using the truth value *unknown*:
 - OR: $(\text{unknown} \text{ or } \text{true}) = \text{true}$
 $(\text{unknown} \text{ or } \text{false}) = \text{unknown}$
 $(\text{unknown} \text{ or } \text{unknown}) = \text{unknown}$
 - AND: $(\text{true} \text{ and } \text{unknown}) = \text{unknown}$
 $(\text{false} \text{ and } \text{unknown}) = \text{false}$
 $(\text{unknown} \text{ and } \text{unknown}) = \text{unknown}$
 - NOT: $(\text{not unknown}) = \text{unknown}$
 - In SQL “*P* is **unknown**” evaluates to true if predicate *P* evaluates to *unknown*
- Result of select predicate is treated as *false* if it evaluates to *unknown*

Equivalent Queries

- There is more than one way to write a query in relational algebra
- Example: Find information about courses taught by instructors in the Physics department with salary greater than 90,000
- Query 1

$$\sigma_{dept_name="Physics" \wedge salary > 90,000} (instructor)$$

- Query 2

$$\sigma_{dept_name="Physics"} (\sigma_{salary > 90,000} (instructor))$$

- The two queries are not identical
- They are, however, equivalent and give the same result on any database

Equivalent Queries

- There is more than one way to write a query in relational algebra
- Example: Find information about courses taught by instructors in the Physics department
- Query 1

$$\sigma_{dept_name="Physics"}(instructor \bowtie_{instructor.ID = teaches.ID} teaches)$$

- Query 2

$$(\sigma_{dept_name="Physics"}(instructor)) \bowtie_{instructor.ID = teaches.ID} teaches$$

- The two queries are not identical
- They are, however, equivalent and give the same result on any database

Notes about Relational Languages

- Each query input is a table or a set of tables
- Each query output is a table
- All data in the output table appears in atleast one of the input tables
- Relational algebra is NOT Turing complete

Summary of Relational Algebra Operators

| Symbol (Name) | Example of Use |
|------------------------------|---|
| σ (selection) | $\sigma_{dept_name="Physics"}(instructor)$ Returns the rows of the input relation that satisfy the predicate |
| Π (projection) | $\Pi_{ID, name, salary}(instructor)$ Output specified attributes from all rows of the input relation Remove duplicate tuples from the output |
| \times (Cartesian product) | $instructor \times teaches$ Output each possible pair of tuples of the two input relations |
| \cup (union) | $\Pi_{course_id}(\sigma_{semester="Fall" \wedge year=2017}(section)) \cup \Pi_{course_id}(\sigma_{semester="Spring" \wedge year=2018}(section))$ Output the union of tuples from the two input relations |
| $-$ (set difference) | $\Pi_{course_id}(\sigma_{semester="Fall" \wedge year=2017}(section)) - \Pi_{course_id}(\sigma_{semester="Spring" \wedge year=2018}(section))$ Output the set difference of tuples from the two input relations |
| \bowtie (natural join) | $instructor \bowtie_{instructor.id = teaches.id} teaches$ Output pairs of rows from the two input relations that have the same value on all attributes that have the same name |

Next Lecture

Introduction to SQL

Thank you for your attention...

Any question?

Contact:

Department of Information Technology, NITK Surathkal, India
6th Floor, Room: 13

Phone: +91-9477678768

E-mail: shrutilipi@nitk.edu.in