



Introduction to SQL

History

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
 - SQL-86
 - SQL-89
 - SQL-92
 - SQL:1999 (language name became Y2K compliant!)
 - SQL:2003
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features
 - Not all examples here may work on your particular system

SQL Parts

- **DML:** Provides the ability to query information from the database and to insert tuples into, delete tuples from, and modify tuples in the database
- **Integrity:** The DDL includes commands for specifying integrity constraints
- **View definition:** The DDL includes commands for defining views
- **Transaction control:** Includes commands for specifying the beginning and ending of transactions
- **Embedded SQL and dynamic SQL:** Define how SQL statements can be embedded within general-purpose programming languages
- **Authorization:** Includes commands for specifying access rights to relations and views

Data Definition Language

- The SQL **data-definition language (DDL)** allows the specification of information about relations, including:
 - The schema for each relation
 - The type of values associated with each attribute
 - The Integrity constraints
 - The set of indices to be maintained for each relation
 - Security and authorization information for each relation
 - The physical storage structure of each relation on disk

Domain Types in SQL

- **char(n):** Fixed length character string, with user-specified length n
- **varchar(n):** Variable length character strings, with user-specified maximum length n
- **int:** Integer (a finite subset of the integers that is machine-dependent)
- **smallint:** Small integer (a machine-dependent subset of the integer domain type)
- **numeric(p , d):** Fixed point number, with user-specified precision of p digits, with d digits to the right of decimal point (ex., **numeric**(3,1), allows 44.5 to be stores exactly, but not 444.5 or 0.32)
- **real, double precision:** Floating point and double-precision floating point numbers, with machine-dependent precision
- **float(n):** Floating point number, with user-specified precision of at least n digits

Create Table Construct

- An SQL relation is defined using the **create table** command:

```
create table r
    (A1 D1, A2 D2, ..., An Dn,
    (integrity-constraint1),
    ...,
    (integrity-constraintk))
```

- r* is the name of the relation
- each *A_i* is an attribute name in the schema of relation *r*
- D_i* is the data type of values in the domain of attribute *A_i*
- Example:

```
create table instructor (
    ID           char(5),
    name        varchar(20),
    dept_name   varchar(20),
    salary      numeric(8, 2))
```

Integrity Constraints in Create Table

- Types of integrity constraints
 - **primary key** (A_1, \dots, A_n)
 - **foreign key** (A_m, \dots, A_n) **references** r
 - **not null**
- SQL prevents any update to the database that violates an integrity constraint
- Example:

```
create table instructor (  
    ID          char(5),  
    name       varchar(20) not null,  
    dept_name varchar(20),  
    salary    numeric(8, 2),  
    primary key (ID),  
    foreign key (dept_name) references department);
```

Relation Definitions

```
create table student (  
    ID          varchar(5),  
    name       varchar(20) not null,  
    dept_name  varchar(20),  
    tot_cred   numeric(3,0),  
    primary key (ID),  
    foreign key (dept_name) references department);
```

```
create table takes (  
    ID          varchar(5),  
    course_id   varchar(8),  
    sec_id      varchar(8),  
    semester    varchar(6),  
    year        numeric(4,0),  
    grade       varchar(2),  
    primary key (ID, course_id, sec_id, semester, year),  
    foreign key (ID) references student,  
    foreign key (course_id, sec_id, semester, year) references section);
```

```
create table course (  
    course_id   varchar(8),  
    title       varchar(50),  
    dept_name   varchar(20),  
    credits     numeric(2,0),  
    primary key (course_id),  
    foreign key (dept_name) references department);
```


Updates to Tables

- **Insert**
 - **insert into** *instructor* **values** ('10211', 'Smith', 'Biology', 66000);
- **Delete**
 - Remove all tuples from the *student* relation
 - **delete from** *student*
- **Drop Table**
 - **drop table** *r*
- **Alter**
 - **alter table** *r* **add** *A D*
 - Where *A* is the name of the attribute to be added to relation *r* and *D* is the domain of *A*
 - All exiting tuples in the relation are assigned *null* as the value for the new attribute
 - **alter table** *r* **drop** *A*
 - Where *A* is the name of an attribute of relation *r*
 - Dropping of attributes not supported by many databases

Basic Query Structure

- A typical SQL query has the form:

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

- A_i represents an attribute
 - R_i represents a relation
 - P is a predicate
- The result of an SQL query is a relation

The *Select* Clause

- The **select** clause lists the attributes desired in the result of a query
 - Corresponds to the projection operation of the relational algebra
- Example: Find the names of all ***instructors***:
select *name* **from** *instructor*
- **Note:** SQL names are case insensitive (i.e., you may use upper- or lower-case letters)
 - E.g., *Name* \equiv *NAME* \equiv *name*
 - Some people use upper case wherever we use bold font

The *Select* Clause

- SQL allows duplicates in relations as well as in query results
- To force the elimination of duplicates, insert the keyword **distinct** after select
- Find the department names of all ***instructors***, and remove duplicates
select distinct dept_name from instructor
- The keyword **all** specifies that duplicates should not be removed
select all dept_name from instructor

<i>dept_name</i>
Comp. Sci.
Finance
Music
Physics
History
Physics
Comp. Sci.
History
Finance
Biology
Comp. Sci.
Elec. Eng.

The *Select* Clause

- An asterisk in the select clause denotes “all attributes”

select * from *instructor*

- An attribute can be a literal with no **from** clause

select '437'

- Results is a table with one column and a single row with value “437”
- Can give the column a name using:

select '437' as AAA

- An attribute can be a literal with **from** clause

select 'A' from *instructor*

- Result is a table with one column and N rows (number of tuples in the *instructors* table), each row with value “A”

The *Select* Clause

- The **select** clause can contain arithmetic expressions involving the operation, +, −, *, and /, and operating on constants or attributes of tuples
- The following query would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12:

```
select ID, name, salary/12 from instructor
```

- Can rename “*salary/12*” using the **as** clause:

```
select ID, name, salary/12 as monthly_salary
```

The *Where* Clause

- The **where** clause specifies conditions that the result must satisfy
 - Corresponds to the selection predicate of the relational algebra
- To find all **instructors** in Comp. Sci. dept


```
select name
from instructor
where dept_name = 'Comp. Sci.'
```
- SQL allows the use of the logical connectives **and**, **or**, and **not**
- The operands of the logical connectives can be expressions involving the comparison operators <, <=, >, >=, =, and <>
- Comparisons can be applied to results of arithmetic expressions
- To find all **instructors** in Comp. Sci. dept with salary > 70000


```
select name
from instructor
where dept_name = 'Comp. Sci.' and salary > 70000
```

<i>name</i>
Katz
Brandt

The *From* Clause

- The **from** clause lists the relations involved in the query
 - Corresponds to the Cartesian product operation of the relational algebra
- Find the Cartesian product ***instructor* × *teaches*** generates every possible instructor – teaches pair, with all attributes from both relations

select * from *instructor*, *teaches*

 - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful when combined with where-clause condition (selection operation in relational algebra)

Examples

- Find the names of all instructors who have taught some course and the *course_id*
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID
- Find the names of all instructors in the Art ***department*** who have taught some course and the *course_id*
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID and
instructor.dept_name = 'Art'

<i>name</i>	<i>course_id</i>
Srinivasan	CS-101
Srinivasan	CS-315
Srinivasan	CS-347
Wu	FIN-201
Mozart	MU-199
Einstein	PHY-101
El Said	HIS-351
Katz	CS-101
Katz	CS-319
Crick	BIO-101
Crick	BIO-301
Brandt	CS-190
Brandt	CS-190
Brandt	CS-319
Kim	EE-181

The *Rename* Operation

- The SQL allows renaming relations and attributes using the **as** clause:

old-name as new-name

- Find the names of all ***instructors*** who have a higher salary than some instructor in 'Comp. Sci'.

select distinct *T.name*

from *instructor as T, instructor as S*

where *T.salary > S.salary and S.dept_name = 'Comp. Sci.'*

- Keyword **as** is optional and may be omitted

instructor as T \equiv *instructor T*

Next Lecture

Introduction to SQL

Thank you for your attention...

Any question?

Contact:

Department of Information Technology, NITK Surathkal, India
6th Floor, Room: 13

Phone: +91-9477678768

E-mail: shrutilipi@nitk.edu.in

Self Join Example

- Relation *emp-super*

<i>person</i>	<i>supervisor</i>
Bob	Alice
Mary	Susan
Alice	David
David	Mary

- Find the supervisor of “Bob”
- Find the supervisor of the supervisor of “Bob”
- Can you find ALL the supervisors (direct and indirect) of “Bob”?