



Normalization

Multivalued Dependencies (MVDs)

- **Persons**(*Man(M)*, *Phone(P)*, *Dogs_Like(D)*)

Person:			Meaning of the tuples
Man (M)	Phone (P)	Dogs_Like (D)	Man M have phones P and likes the dog D
M1	P1/P2	D1/D2	M1 have phones P1 and P2, and likes the dogs D1 and D2
M2	P3	D2	M2 have phones P3, and likes the dog D2
Key: MPD			

- There are no non-trivial FDs because all tributes are combined forming *candidate key*, i.e., MDP
- In the above relation, two multivalued dependency exists:
 - $Man \twoheadrightarrow Phones$
 - $Man \twoheadrightarrow Dogs_Like$
- A Man's phone are independent of the dogs they like
- But, after converting the above relation in *single valued attribute*, each of a man's phones appears with each of the dogs they like in all combinations

Post 1NF normalization		
Man (M)	Phone (P)	Dogs_Like (D)
M1	P1	D1
M1	P2	D2
M2	P3	D2
M2	P3	D2
M1	P1	D2
M1	P2	D1

Multivalued Dependencies (MVDs)

- If two or more independent relations are kept in a single relation, then *multivalued dependency* is possible
- For example, let there are two relations:
 - **Student**(SID, Sname) where $SID \rightarrow Sname$
 - **Course**(CID, Cname) where $CID \rightarrow Cname$
- There is no relation defined between Student and Course
- If we kept them in a single relations, named **Student_Course**, then MVD will exists because of *m:n cardinality*
- If two or more *MVDs* exists in a relation, then while converting into *SVAs*, *MVDs* exits

Student:	
SID	Sname
S1	A
S2	B

Course:	
CID	Cname
C1	C
C2	B

SID	Sname	CID	Cname
S1	A	C1	C
S1	A	C2	B
S2	B	C1	C
S2	B	C2	B

Multivalued Dependencies (MVDs)

- Suppose we record names of children, and phone numbers for instructors:
 - *inst_child*(*ID*, *child_name*)
 - *inst_phone*(*ID*, *phone_number*)
- If we were to combine these schemas to get
 - *inst_info*(*ID*, *child_name*, *phone_number*)
 - Example data:
 - (99999, David, 512-555-1234)
 - (99999, David, 512-555-4321)
 - (99999, William, 512-555-1234)
 - (99999, William, 512-555-4321)
- This relation is in BCNF
 - Why?

Multivalued Dependencies (MVDs)

- Let R be a relation schema and let $\alpha \subseteq R$ and $\beta \subseteq R$. The **multivalued dependency**
 $\alpha \twoheadrightarrow \beta$

holds on R if in any legal relation $r(R)$, for all pairs for tuples t_1 and t_2 in r such that $t_1[\alpha] = t_2[\alpha]$, there exist tuples t_3 and t_4 in r such that:

$$t_1[\alpha] = t_2[\alpha] = t_3[\alpha] = t_4[\alpha]$$

$$t_3[\beta] = t_1[\beta]$$

$$t_3[R - \beta] = t_2[R - \beta]$$

$$t_4[\beta] = t_2[\beta]$$

$$t_4[R - \beta] = t_1[R - \beta]$$

Example:

A relation of university courses, the books recommended for the course, and the lectures who will be teaching the course:

$course \twoheadrightarrow book$

$course \twoheadrightarrow lecturer$

Course	Book	Lecturer	Tuples
AHA	Silberschatz	John D	t1
AHA	Nederpelt	William M	t2
AHA	Silberschatz	William M	t3
AHA	Nederpelt	John D	t4
AHA	Silberschatz	Christian G	
AHA	Nederpelt	Christian G	
OSO	Silberschatz	John D	
OSO	Silberschatz	William M	

MVD: Tabular Representation

- Tabular representation of $\alpha \twoheadrightarrow \beta$

	α	β	$R - \alpha - \beta$
t_1	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$a_{j+1} \dots a_n$
t_2	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$b_{j+1} \dots b_n$
t_3	$a_1 \dots a_i$	$a_{i+1} \dots a_j$	$b_{j+1} \dots b_n$
t_4	$a_1 \dots a_i$	$b_{i+1} \dots b_j$	$a_{j+1} \dots a_n$

MVD

- Let R be a relation schema with a set of attributes that are partitioned into 3 nonempty subsets

Y, Z, W

- We say that $Y \twoheadrightarrow Z$ (Y **multidetermines** Z) if and only if for all possible relations $r(R)$

$\langle y_1, z_1, w_1 \rangle \in r$ and $\langle y_1, z_2, w_2 \rangle \in r$

- Then

$\langle y_1, z_1, w_2 \rangle \in r$ and $\langle y_1, z_2, w_1 \rangle \in r$

- Note that since the behavior of Z and W are identical it follows that

$Y \twoheadrightarrow Z$ if $Y \twoheadrightarrow W$

Example

- In our example:

$$ID \twoheadrightarrow child_name$$
$$ID \twoheadrightarrow phone_number$$

- The above formal definition is supposed to formalize the notion that given a particular value of Y (ID) it has associated with it a set of values of Z ($child_name$) and a set of values of W ($phone_number$), and these two sets are in some sense independent of each other
- **Note:**
 - If $Y \rightarrow Z$ then $Y \twoheadrightarrow Z$
 - Indeed we have (in above notation) $Z_1 = Z_2$The claim follows

Use of Multivalued Dependencies

- We use multivalued dependencies in two ways:
 - To test relations to **determine** whether they are legal under a given set of functional and multivalued dependencies
 - To specify **constraints** on the set of legal relations
 - We shall concern ourselves *only* with relations that satisfy a given set of functional and multivalued dependencies
- If a relation r fails to satisfy a given multivalued dependency, we can construct a relations r' that does satisfy the multivalued dependency by adding tuples to r

Theory of MVDs

Name		Rule
C-	Complementation	: If $X \twoheadrightarrow Y$, then $X \twoheadrightarrow \{R - (X \cup Y)\}$
A-	Augmentation	: If $X \twoheadrightarrow Y$ and $W \supseteq Z$, then $WX \twoheadrightarrow YZ$
T-	Transitivity	: If $X \twoheadrightarrow Y$ and $Y \twoheadrightarrow Z$, then $X \twoheadrightarrow (Z - Y)$
	Replication	: If $X \rightarrow Y$ and $X \twoheadrightarrow Y$ but the reverse is not true
	Coalescence	: If $X \twoheadrightarrow Y$ and there is a W such that $W \cap Y$ is empty, $W \rightarrow Z$, and $Y \supseteq Z$, then $X \rightarrow Z$

- A MVD $X \twoheadrightarrow Y$ in R is called a trivial MVD is:
 - Y is a subset of X ($X \supseteq Y$) or
 - $X \cup Y = R$
 - Otherwise, it is a non-trivial MVD and we have to repeat values redundancy in the tuples

Theory of MVDs

- From the definition of multivalued dependency, we can derive the following rule:
 - If $\alpha \rightarrow \beta$, then $\alpha \twoheadrightarrow \beta$

That is, every functional dependency is also a multivalued dependency

- The **closure** D^+ of D is the set of all functional and multivalued dependencies logically implied by D
 - We can compute D^+ from D , using the formal definitions of functional dependencies and multivalued dependencies
 - We can manage with such reasoning for very simple multivalued dependencies, which seem to be most common in practice
 - For complex dependencies, it is better to reason about sets of dependencies using a system of inference rules

Fourth Normal Form

- A relation schema R is in **4NF** with respect to a set D of functional and multivalued dependencies if for all multivalued dependencies in D^+ of the form $\alpha \twoheadrightarrow \beta$, where $\alpha \subseteq R$ and $\beta \subseteq R$, at least one of the following hold:
 - $\alpha \twoheadrightarrow \beta$ is trivial (i.e., $\beta \subseteq \alpha$ or $\alpha \cup \beta = R$)
 - α is a *superkey* for schema R
- If a relation is in 4NF, it is in BCNF

Restriction of Multivalued Dependencies

- The restriction of D to R_i is the set D_i consisting of
 - All functional dependencies in D^+ that include only attributes of R_i
 - All multivalued dependencies of the form

$$\alpha \twoheadrightarrow (\beta \cap R_i)$$

where $\alpha \subseteq R_i$ and $\alpha \twoheadrightarrow \beta$ is in D^+

4NF Decomposition Algorithm

- For all dependencies $A \twoheadrightarrow B$ in D^+ , check if A is a *superkey*
 - By using attribute closure
- If not, then
 - Choose a dependency in F^+ that breaks the 4NF rules, say $A \twoheadrightarrow B$
 - Create $R_1 = AB$
 - Create $R_2 = A (R - (B - A))$
 - Note that: $R_1 \cap R_2 = A$ and $A \twoheadrightarrow AB (=R_1)$, so this is lossless decomposition
- Repeat for R_1 and R_2
 - By defining D_1^+ to be all dependencies in F that contain only attributes in R_1
 - Similarly D_2^+

4NF Decomposition Algorithm

```

result := {R};
done := false;
compute  $D^+$ ;
Let  $D_i$  denote the restriction of  $D^+$  to  $R_i$ 
while (not done)
    if (there is a schema  $R_i$  in result that is not in 4NF) then
        begin
            let  $\alpha \twoheadrightarrow \beta$  be a nontrivial multivalued dependency that holdson  $R_i$  such that  $\alpha \rightarrow R_i$ 
                is not in  $D_i$ , and  $\alpha \cap \beta = \emptyset$ ;
            result := (result -  $R_i$ )  $\cup$  ( $R_i$  -  $\beta$ )  $\cup$  ( $\alpha, \beta$ );
        end
    else done := true;

```

Note: Each R_i is in 4NF, and decomposition is lossless-join

Example of 4NF Decomposition

- Example:
- **Persons_Modify**(*Man*(*M*), *Phone*(*P*), *Dogs_Like*(*D*), *Address*(*A*))
- *FDs*:
 - FD1: *Man* \twoheadrightarrow *Phones*
 - FD2: *Man* \twoheadrightarrow *Dogs_Like*
 - FD3: *Man* \rightarrow *Address*
- Key: *MPD*
- All dependencies violate 4NF



Man (M)	Phone (P)	Dogs_Like (D)	Address (A)
M1	P1	D1	49-ABC, Bhiwani (HR.)
M1	P2	D2	49-ABC, Bhiwani (HR.)
M2	P3	D2	36-XYZ, Rohtak (HR.)
M1	P1	D2	49-ABC, Bhiwani (HR.)
M1	P2	D1	49-ABC, Bhiwani (HR.)

- In the above relations for both the MVD's, "**X**" is **Man**, which is again not the *superkey*, but as $X \cup Y = R$, i.e., (*Man* & *Phone*) together make the relations
- So, the above MVDs are trivial and in FD3, *Address* is functionally dependent on *Man*, where *Man* is the key in **Person_Address**, hence all the three relations are in 4NF

Example

- $R = (A, B, C, G, H, I)$
- $F = \{ A \twoheadrightarrow B$
 $B \twoheadrightarrow HI$
 $CG \twoheadrightarrow H \}$
- R is not in 4NF since $A \twoheadrightarrow B$ and A is not a superkey for R Decomposition
 - a) $R_1 = (A, B)$ (R_1 is in 4NF)
 - b) $R_2 = (A, C, G, H, I)$ (R_2 is not in 4NF, decompose into R_3 and R_4)
 - c) $R_3 = (C, G, H)$ (R_3 is in 4NF)
 - d) $R_4 = (A, C, G, I)$ (R_4 is not in 4NF, decompose into R_5 and R_6)
- $A \twoheadrightarrow B$ and $B \twoheadrightarrow HI \Rightarrow A \twoheadrightarrow HI$, (MVD transitivity)
- And hence $A \twoheadrightarrow I$ (MVD restriction to R_4)
 - e) $R_5 = (A, I)$ (R_5 is in 4NF)
 - f) $R_6 = (A, C, G)$ (R_6 is in 4NF)

Design Goals

- Goal for a relational database design is:
 - BCNF/ 4NF
 - Lossless join
 - Dependency preservation
- If we cannot achieve this, we accept one of the following:
 - Lack of dependency preservation
 - Redundancy due to use of 3NF
- Interestingly, SQL does not provide a direct way of specifying functional dependencies other than superkeys
- Can specify FDs using assertions, but they are expensive to test, (and currently not supported by any of the widely used databases!)
- Even if we had a dependency preserving decomposition, using SQL we would not be able to efficiently test a functional dependency whose left hand side is not a key

Further Normal Forms

- **Further NFs**
 - Elementary key normal form (EKNF)
 - Essential tuple normal form (ETNF)
 - Join dependencies and Fifth normal form (5NF)
 - Sixth normal form (6NF)
 - Domain/key normal form (DKNF)
- **Join dependencies** generalize multivalued dependencies
 - Lead to **project-join normal form (PJNF)** (also called **fifth normal form**)
- A class of even more general constraints, leads to a normal form called **domain-key normal form**
- Problem with these generalized constraints: are hard to reason with, and no set of sound and complete set of inference rules exists
- Hence rarely used

Overall Database Design Process

We have assumed schema R is given

- R could have been generated when converting ER diagram to a set of tables
- R could have been a single relation containing *all* attributes that are of interest (called **universal relation**)
- Normalization breaks R into smaller relations
- R could have been the result of some ad hoc design of relations, which we then test/convert to normal form

ER Model and Normalization

- When an ER diagram is carefully designed, identifying all entities correctly, the tables generated from the ER diagram should not need further normalization
- However, in a real (imperfect) design, there can be functional dependencies from non-key attributes of an entity to other attributes of the entity
 - Example: An *employee* entity with
 - Attributes: *department_name* and *building*
 - Functional dependency: *department_name* → *building*
- Good design would have made department an entity
- Functional dependencies from non-key attributes of a relationship set possible, but rare --- most relationships are binary

Denormalization for Performance

- May want to use non-normalized schema for performance
- For example, displaying *prereqs* along with *course_id*, and *title* requires join of *course* with *prereq*
 - **Course**(*course_id*, *title*, ...)
 - **Prerequisite**(*course_id*, *prereq*)
- Alternative 1: Use denormalized relation containing attributes of *course* as well as *prereq* with all above attributes: **Course**(*course_id*, *title*, *prereq*)
 - Faster lookup
 - Extra space and extra execution time for updates
 - Extra coding work for programmer and possibility of error in extra code
- Alternative 2: Use a materialized view defined a *course prereq*
 - **Course** ⋈ **Prerequisite**
 - Benefits and drawbacks same as above, except no extra coding work for programmer and avoids possible errors

Other Design Issues

- Some aspects of database design are not caught by normalization
- Examples of bad database design, to be avoided:

Instead of *earnings* (*company_id*, *year*, *amount*), use

- *earnings_2004*, *earnings_2005*, *earnings_2006*, etc., all on the schema (*company_id*, *earnings*)
 - Above are in BCNF, but make querying across years difficult and needs new table each year
- *company_year* (*company_id*, *earnings_2004*, *earnings_2005*, *earnings_2006*)
 - Also in BCNF, but also makes querying across years difficult and requires new attribute each year
 - Is an example of a **crosstab**, where values for one attribute become column names
 - Used in spreadsheets, and in data analysis tools

Modeling Temporal Data

- **Temporal data** have an association time interval during which the data are *valid*
- A **snapshot** is the value of the data at a particular point in time
- Several proposals to extend ER model by adding valid time to
 - Attributes, e.g., Address of an instructor at different points in time
 - Entities, e.g., Time duration when a student entity exists
 - Relationships, e.g., Time during which an instructor was associated with a student as an advisor
- But no accepted standard
- Adding a temporal component results in functional dependencies like

$$ID \rightarrow \text{street, city}$$

not holding, because the address varies over time

- A **temporal functional dependency** $X \rightarrow Y$ holds on schema R if the functional dependency $X \rightarrow Y$ holds on all snapshots for all legal instances $r(R)$

Modeling Temporal Data

- In practice, database designers may add start and end time attributes to relations
 - E.g., *course(course_id, course_title)* is replaced by
course(course_id, course_title, start, end)
 - Constraint: No two tuples can have overlapping valid times
 - Hard to enforce efficiently
- Foreign key references may be to current version of data, or to data at a point in time
 - E.g., Student transcript should refer to course information at the time the course was taken

Extra

Correctness of 3NF Decomposition Algorithm

- 3NF decomposition algorithm is dependency preserving (since there is a relation for every FD in F_c)
- Decomposition is lossless
 - A candidate key (C) is in one of the relations R_i in decomposition
 - Closure of candidate key under F_c must contain all attributes in R
 - Follow the steps of attribute closure algorithm to show there is only one tuple in the join result for each tuple in R_i

Correctness of 3NF Decomposition Algorithm

- Claim: If a relation R_i is in the decomposition generated by the above algorithm, then R_i satisfies 3NF
- Proof:
 - Let R_i be generated from the dependency $\alpha \rightarrow \beta$
 - Let $\gamma \rightarrow B$ be any non-trivial functional dependency on R_i
(We need only consider FDs whose right-hand side is a single attribute)
 - Now, B can be in either β or α but not in both
 - Consider each case separately

Correctness of 3NF Decomposition Algorithm

- Case 1: If B in β :
 - If γ is a superkey, the 2nd condition of 3NF is satisfied
 - Otherwise α must contain some attribute not in γ
 - Since $\gamma \rightarrow B$ is in F^+ it must be derivable from F_c , by using attribute closure on γ
 - Attribute closure not have used $\alpha \rightarrow \beta$
 - If it had been used, α must be contained in the attribute closure of γ , which is not possible, since we assumed γ is not a superkey
 - Now, using $\alpha \rightarrow (\beta - \{B\})$ and $\gamma \rightarrow B$, we can derive $\alpha \rightarrow B$
(since $\gamma \subseteq \alpha$, and $B \notin \gamma$ since $\gamma \rightarrow B$ is non-trivial)
 - Then, B is extraneous in the right-hand side of $\alpha \rightarrow \beta$; which is not possible since $\alpha \rightarrow \beta$ is in F_c
 - Thus, if B is in β then γ must be a superkey, and the second condition of 3NF must be satisfied

Correctness of 3NF Decomposition Algorithm

- Case 2: B is in α
 - Since α is a candidate key, the third alternative in the definition of 3NF is trivially satisfied
 - In fact, we cannot show that γ is a superkey
 - This shows exactly why the third alternative is present in the definition of 3NF

Next Lecture

Transactions

Thank you for your attention...

Any question?

Contact:

Department of Information Technology, NITK Surathkal, India
6th Floor, Room: 13

Phone: +91-9477678768

E-mail: shrutilipi@nitk.edu.in