



## Transactions

# Serializability

- **Basic Assumption:** Each transaction preserves database consistency
- Thus, serial execution of a set of transactions preserves database consistency
- A (possibly concurrent) **schedule is serializable if it is equivalent to a serial schedule**
- Different forms of schedule equivalence give rise to the notions of:
  - **Conflict serializability**
  - **View serializability**

# Simplified View of Transactions

- We ignore operations other than **read** and **write** instructions
  - Other operations happen in memory (are temporary in nature) and (mostly) do not affect the state of the database
  - This is simplifying assumptions for analysis
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes
- Our simplified schedules consist of only **read** and **write** instructions

# Conflicting Instructions

- Instructions  $I_i$  and  $I_j$  of transactions  $T_i$  and  $T_j$  respectively
- Instructions  $I_i$  and  $I_j$  **conflict** if and only if there exists some item  $Q$  accessed by both  $I_i$  and  $I_j$ , and at least one of these instructions wrote  $Q$ 
  - 1.  $I_i = \text{read}(Q), I_j = \text{read}(Q) \rightarrow I_i$  and  $I_j$  don't conflict
  - 2.  $I_i = \text{read}(Q), I_j = \text{write}(Q) \rightarrow$  They conflict
  - 3.  $I_i = \text{write}(Q), I_j = \text{read}(Q) \rightarrow$  They conflict
  - 4.  $I_i = \text{write}(Q), I_j = \text{write}(Q) \rightarrow$  They conflict
- Intuitively, a conflict between  $I_i$  and  $I_j$  forces a (logical) temporal order between them
  - If  $I_i$  and  $I_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule

# Conflict Serializability

- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**
- We say that a schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule

# Conflict Serializability

- Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions
  - Swap  $T_1.\text{read}(B)$  and  $T_2.\text{write}(A)$
  - Swap  $T_1.\text{read}(B)$  and  $T_2.\text{read}(A)$
  - Swap  $T_1.\text{write}(B)$  and  $T_2.\text{write}(A)$
  - Swap  $T_1.\text{write}(B)$  and  $T_2.\text{read}(A)$
- Therefore Schedule 3 is conflict serializable

These swaps do not conflict as they work with different items ( $A$  &  $B$ ) in different transaction

$T_1$	$T_2$
read (A) write (A)	
	read (A) write (A)
read (B) write (B)	
	read (B) write (B)

Schedule 3

$T_1$	$T_2$
read (A) write (A)	
	read (A)
read (B)	write (A)
write (B)	read (B) write (B)

Schedule 5

$T_1$	$T_2$
read (A) write (A) read (B) write (B)	
	read (A) write (A) read (B) write (B)

Schedule 6

# Conflict Serializability

- Example of a schedule that is not conflict serializable:

$T_3$	$T_4$
read (Q)	
write (Q)	write (Q)

- We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$

# Example: Bad Schedule

- Consider two transactions:

## Transaction 1

**UPDATE** accounts

**SET** balance = balance - 100

**WHERE** acct\_id = 31414

## Transaction 2

**UPDATE** accounts

**SET** balance = balance \* 1.005

- In terms of read-write, we can write as follows:
  - Transaction 1:  $r_1(A)$ ,  $w_1(A)$  // A is the balance for acct\_id = 31414
  - Transaction 2:  $r_2(A)$ ,  $w_2(A)$ ,  $r_2(B)$ ,  $w_2(B)$  // B is the balance of other accounts
- Consider schedule S
  - Schedule S:  $r_1(A)$ ,  $r_2(A)$ ,  $w_1(A)$ ,  $w_2(A)$ ,  $r_2(B)$ ,  $w_2(B)$
  - Suppose: Account A starts with \$200, and account B starts with \$100
  - Schedule S is very bad! (At least, it's bad if you're the bank!) We withdrew \$100 from account A, but somehow the database has recorded that our account now holds \$201

	A	B
(initial:)	200.00	100.00
$r_1(A)$ :		
$r_2(A)$ :		
$w_1(A)$ :	100.00	
$w_2(A)$ :	201.00	
$r_2(B)$ :		
$w_2(B)$ :		100.50

Schedule S



# Example: Bad Schedule

- Ideal schedule is **serial**: ( $A = \$200$ ,  $B = \$100$ )
  - Serial schedule 1:  $r_1(A)$ ,  $w_1(A)$ ,  $r_2(A)$ ,  $w_2(A)$ ,  $r_2(B)$ ,  $w_2(B)$  //  $A = 100.50$ ,  $B = 100.50$
  - Serial schedule 2:  $r_2(A)$ ,  $w_2(A)$ ,  $r_2(B)$ ,  $w_2(B)$ ,  $r_1(A)$ ,  $w_1(A)$  //  $A = 101.00$ ,  $B = 100.50$

- We call a schedule **serializable** if it has the same effect as some serial schedule regardless of the specific information in the database

- As an example, consider Schedule  $T$ , which has swapped the third and fourth operations from  $S$ :

- Schedule  $S$ :  $r_1(A)$ ,  $r_2(A)$ ,  $w_1(A)$ ,  $w_2(A)$ ,  $r_2(B)$ ,  $w_2(B)$
- Schedule  $T$ :  $r_1(A)$ ,  $r_2(A)$ ,  $w_2(A)$ ,  $w_1(A)$ ,  $r_2(B)$ ,  $w_2(B)$

- Looking just at the first example, we see that the outcome is the same as the serial schedule where the withdrawal happens first and then the interest is credited

- But that's just a peculiarity of the data, as revealed by the second example, where the final value of  $A$  can't be the consequence of either of the possible serial schedules

$A$ is \$100 initially			$A$ is \$200 initially		
	$A$	$B$		$A$	$B$
(initial:)	100.00	100.00	(initial:)	200.00	100.00
$r_1(A)$ :			$r_1(A)$ :		
$r_2(A)$ :			$r_2(A)$ :		
$w_2(A)$ :	100.50		$w_2(A)$ :	201.00	
$w_1(A)$ :	0.00		$w_1(A)$ :	100.00	
$r_2(B)$ :			$r_2(B)$ :		
$w_2(B)$ :		100.50	$w_2(B)$ :		100.50

# Example: Good Schedule

- What's a non-serial example of a serializable schedule?
  - We could credit interest to  $A$  first, then withdraw the money, then credit interest to  $B$ :
  - Schedule  $U$ :  $r_2(A), w_2(A), r_1(A), w_1(A), r_2(B), w_2(B)$
- Schedule  $U$  is conflict-equivalent to Serial Schedule 2, as shown by the following series of swaps:

Schedule  $U$ :

swap  $w_1(A)$  and  $r_2(B)$ :

swap  $w_1(A)$  and  $w_2(B)$ :

swap  $r_1(A)$  and  $r_2(B)$ :

swap  $r_1(A)$  and  $w_2(B)$ :

$r_2(A), w_2(A), r_1(A), w_1(A), r_2(B), w_2(B)$

$r_2(A), w_2(A), r_1(A), r_2(B), w_1(A), w_2(B)$

$r_2(A), w_2(A), r_1(A), r_2(B), w_2(B), w_1(A)$

$r_2(A), w_2(A), r_2(B), r_1(A), w_2(B), w_1(A)$

$r_2(A), w_2(A), r_2(B), w_2(B), r_1(A), w_1(A)$


= Schedule 2

# Serializability

- Are all serializable schedules conflict-serializable? No
- Consider the following schedule for a set of three transactions:
  - $w_1(A), w_2(A), w_2(B), w_1(B), w_3(B)$
- We can perform no swaps to this:
  - The first two operations are both on  $A$  and at least one is a write
  - The second and third operations are by the same transaction
  - The third and fourth are both on  $B$  at at least one is a write, and
  - So are the fourth and fifth
  - So this schedule is not conflict-equivalent to anything else- and certainly not any serial schedules
- However, since nobody ever reads the values written by the  $w_1(A)$ ,  $w_2(B)$ , and  $w_1(B)$  operations, the schedule has the same outcome as the serial schedule:
  - $w_1(A), w_1(B), w_2(A), w_2(B), w_3(B)$




# Precedence Graph

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- Precedence graph**
  - A directed graph where the vertices are the transactions (names)
- We draw an arc from  $T_i$  to  $T_j$  if the two transaction conflict, and  $T_i$  accessed the data item on which the conflict arose earlier
- We may label the arc by the item that was accessed
- Example of a precedence graph



# Test for Conflict Serializability

- A schedule is conflict serializable if and only if its precedence graph is acyclic
- Cycle-detection algorithms exist which take order  $n^2$  time, where  $n$  is the number of vertices in the graph
  - Better algorithms take order  $n + e$  where  $e$  is the number of edges
- If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph
  - This is a linear order consistent with the partial order of the graph
  - For example, a serializability order for Schedule (a) would be one of either (b) or (c)




# Test for Conflict Serializability

- Build a directed graph, with a vertex for each transaction
- Go through each operation of the schedule
  - If the operation is of the form  $w_i(X)$ , find each subsequent operation in the schedule also operating on the same data element  $X$  by a different transaction, that is, anything of the form  $r_j(X)$  or  $w_j(X)$ 
    - For each such subsequent operation, add a directed edge in the graph from  $T_i$  to  $T_j$
  - If the operation is of the form  $r_i(X)$ , find each subsequent *write* to the same data element  $X$  by a different transaction, that is, anything of the form  $w_j(X)$ 
    - For each such subsequent write, add a directed edge in the graph from  $T_i$  to  $T_j$
- The schedule is conflict-serializable if and only if the resulting directed graph is acyclic

# Test for Conflict Serializability

- Consider the following schedule:
  - $w_1(A), r_2(A), w_1(B), w_3(C), r_2(C), r_4(B), w_2(D), w_4(E), r_5(D), w_5(E)$
- We start with an empty graph with five vertices labeled  $T_1, T_2, T_3, T_4, T_5$
- We go through each operation in the schedule:
  - $w_1(A)$ :  $A$  is subsequently read by  $T_2$ , so add edge  $T_1 \rightarrow T_2$
  - $r_2(A)$ : no subsequent writes to  $A$ , so no new edges
  - $w_1(B)$ :  $B$  is subsequently read by  $T_4$ , so add edge  $T_1 \rightarrow T_4$
  - $w_3(C)$ :  $C$  is subsequently read by  $T_2$ , so add edge  $T_3 \rightarrow T_2$
  - $r_2(C)$ : no subsequent writes to  $C$ , so no new edges
  - $r_4(B)$ : no subsequent writes to  $B$ , so no new edges
  - $w_2(D)$ :  $C$  is subsequently read by  $T_2$ , so add edge  $T_3 \rightarrow T_2$
  - $w_4(E)$ :  $E$  is subsequently written by  $T_5$ , so add edge  $T_4 \rightarrow T_5$
  - $r_5(D)$ : no subsequent writes to  $D$ , so no new edges
  - $w_5(E)$ : no subsequent operations on  $E$ , so no new edges

- We end up with the following directed graph



- This graph has no cycles, so the original schedule must be serializable
- Moreover, since one way to topologically sort the graph is  $T_3 - T_1 - T_4 - T_2 - T_5$ , one serial schedule that is conflict-equivalent is:
  - $w_3(C), w_1(A), w_1(B), r_4(B), w_4(E), r_2(A), r_2(C), w_2(D), r_5(D), w_5(E)$

# Next Lecture

## Transactions



# Thank you for your attention...

Any question?

**Contact:**

Department of Information Technology, NITK Surathkal, India  
6<sup>th</sup> Floor, Room: 13

**Phone:** +91-9477678768

**E-mail:** [shrutilipi@nitk.edu.in](mailto:shrutilipi@nitk.edu.in)