



Concurrency Control

Deadlock Handling

- System is deadlocked if there is a set of transactions such that every transaction in the set is waiting for another transaction in the set
- **Deadlock prevention** protocols ensure that the system will *never* enter into a deadlock state
- Some prevention strategies:
 - Require that each transaction locks all its data items before it begins execution (pre-declaration)
 - Impose partial ordering of all data items and require that a transaction can lock data items only in the order specified by the partial order (graph-based protocol)

Deadlock Prevention

- Following schemes use transaction timestamps for the sake of deadlock prevention alone
- **Wait-die** scheme: Non-preemptive
 - Older transaction may wait for younger one to release data item (older means smaller timestamp)
 - Younger transactions never wait for older ones; they are rolled back instead
 - A transaction may die several times before acquiring needed data item
- **Wound-wait** scheme: Preemptive
 - Older transaction *wounds* (forces rollback) of younger transaction instead of waiting for it
 - Younger transactions may wait for older ones
 - May be fewer rollbacks than *wait-die* scheme


Deadlock Prevention

- Both in *wait-die* and in *wound-waits* schemes, a rolled back transactions is restarted with its original timestamp
- Older transactions thus have precedence over newer ones, and starvation is hence avoided
- **Timeout-Based Schemes:**
 - A transaction waits for a lock only for a specified amount of time
 - If the lock has not been granted within that time, the transaction is rolled back and restarted
 - Thus, deadlocks are not possible
 - Simple to implement; but starvation is possible
 - Also difficult to determine good value of the timeout interval


Deadlock Detection

- Deadlocks can be described as a *wait-for graph*, which consists of a pair $G = (V, E)$
 - V is a set of vertices (all the transactions in the system)
 - E is a set of edges; each element is an ordered pair $T_i \rightarrow T_j$
- If $T_i \rightarrow T_j$ is in E , then there is a directed edge from T_i to T_j , implying that T_i is waiting for T_j to release a data item
- When T_i requests a data item currently being held by T_j , then the edge $T_i \rightarrow T_j$ is inserted in the wait-for graph
- This edge is removed only when T_j is no longer holding a data item needed by T_i
- The system is in a deadlock state if and only if the wait-for graph has a cycle
- Must invoke a deadlock-detection algorithm periodically to look for cycles

Deadlock Detection: Example



Wait-for graph without a cycle



Wait-for graph with a cycle

Deadlock Recovery

- When deadlock is detected :
 - Some transaction will have to rolled back (made a **victim**) to break deadlock cycle
 - Select that transaction as victim that will incur minimum cost
 - Rollback: Determine how far to roll back transaction
 - **Total rollback:** Abort the transaction and then restart it
 - More effective to roll back transaction only as far as necessary to break deadlock
 - Starvation happens if same transaction is always chosen as victim
 - Include the number of rollbacks in the cost factor to avoid starvation

Timestamp-Based Protocols

- Each transaction is issued a timestamp when it enters the system
- If an old transaction T_i has time-stamp $TS(T_i)$, a new transaction T_j is assigned time-stamp $TS(T_j)$ such that $TS(T_i) < TS(T_j)$
- The protocol manages concurrent execution such that the time-stamps determine the serializability order
- In order to assure such behavior, the protocol maintains for each data Q two timestamp values:
 - **W-timestamp**(Q) is the largest time-stamp of any transaction that executed **write**(Q) successfully
 - **R-timestamp**(Q) is the largest time-stamp of any transaction that executed **read**(Q) successfully

Timestamp-Based Protocols

- The timestamp ordering protocol ensures that any conflicting **read** and **write** operations are executed in timestamp order
- Suppose a transaction T_i issues a **read**(Q)
 - If $TS(T_i) \leq \mathbf{W}$ -timestamp(Q), then T_i needs to read a value of Q that was already overwritten
 - Hence, the **read** operation is rejected, and T_i is rolled back
 - If $TS(T_i) \geq \mathbf{W}$ -timestamp(Q), then the **read** operation is executed, and **R**-timestamp(Q) is set to $\max(\mathbf{R}$ -timestamp(Q), $TS(T_i))$

Timestamp-Based Protocols

- Suppose that transaction T_i issues **write**(Q)
 - If $TS(T_i) < \mathbf{R}$ -timestamp(Q), then the value of Q that T_i is producing was needed previously, and the system assumed that that value would never be produced
 - Hence, the **write** operation is rejected, and T_i is rolled back
 - If $TS(T_i) < \mathbf{W}$ -timestamp(Q), then T_i is attempting to write an obsolete value of Q
 - Hence, this **write** operation is rejected, and T_i is rolled back
 - Otherwise, the **write** operation is executed, and \mathbf{W} -timestamp(Q) is set to $TS(T_i)$


Example Use of the Protocol

- A partial schedule for several data items for transactions with timestamps 1, 2, 3, 4, 5

T_1	T_2	T_3	T_4	T_5
				read (X)
read (Y)	read (Y)	write (Y) write (Z)		
				read (Z)
	read (Z) abort			
read (X)		write (W) abort	read (W)	
				write (Y) write (Z)

Correctness of Timestamp-Ordering Protocol

- The timestamp-ordering protocol guarantees serializability since all the arcs in the precedence graph are of the form:



- Thus, there will be no cycles in the precedence graph
- Timestamp protocol ensures freedom from deadlock as no transaction ever waits
- But the schedule may not be cascade-free, and may not even be recoverable

Next Lecture

Recovery System

Thank you for your attention...

Any question?

Contact:

Department of Information Technology, NITK Surathkal, India
6th Floor, Room: 13

Phone: +91-9477678768

E-mail: shrutilipi@nitk.edu.in