



## **Basic Concepts of Indexing**

# Search Records

- Consider a table: **Faculty**(Name, Phone)

Index on "Name"		Table "Faculty"			Index on "Phone"	
Name	Pointer	Rec #	Name	Phone	Pointer	Phone
Anupam Basu	2	1	Partha Pratim Das	81998	6	81664
Pabitra Mitra	6	2	Anupam Basu	82404	1	81998
Partha Pratim Das	1	3	Ranjan Sen	84624	2	82404
Prabir Kumar Biswas	7	4	SudeshnaSarkar	82432	4	82432
Rajib Mall	5	5	Rajib Mall	83668	5	83668
Ranjan Sen	3	6	Pabitra Mitra	81664	3	84624
SudeshnaSarkar	4	7	Prabir Kumar Biswas	84772	7	84772

- How to search on Name?
  - Get the phone number for 'Pabitra Mitra'
  - Use "Name" Index – sorted on 'Name', search 'Pabitra Mitra' and navigate on pointer (rec #)
- How to search on Phone?
  - Get the name of the faculty having phone number = 84772
  - Use "Phone" Index – sorted on 'Phone', search '84772' and navigate on pointer (rec #)
- We can keep the records sorted on 'Name' or on 'Phone' (called the primary index), but not on both

# Basic Concepts

- Indexing mechanisms used to speed up access to desired data
- For example:
  - Name in a faculty table
  - Author catalog in library
- **Search Key** attribute to set of attributes used to look up records in a file
- An **index file** consists of records (called **index entries**) of the form



- Index files are typically much smaller than the original file
- Two basic kinds of indices:
  - **Ordered indices:** Search keys are stored in sorted order
  - **Hash indices:** Search keys are distributed uniformly across “buckets” using a “hash function”

# Index Evaluation Metrics

- Access types supported efficiently
- For example:
  - Records with a specified value in the attribute
  - Records with an attribute value falling in a specified range of values
- Access time
- Insertion time
- Deletion time
- Space overhead

# Ordered Indices

- In an **ordered index**, index entries are stored sorted on the search key value
  - For example, author catalog in library
- **Primary index**: In a sequentially ordered file, the index whose search key specifies the sequential order of the file
  - Also called **clustering index**
  - The search key of a primary index is usually but not necessarily the primary key
- **Secondary index**: An index whose search key specifies an order different from the sequential order of the file
  - Also called **nonclustering index**
- **Index-sequential file**: Ordered sequential file with a primary index

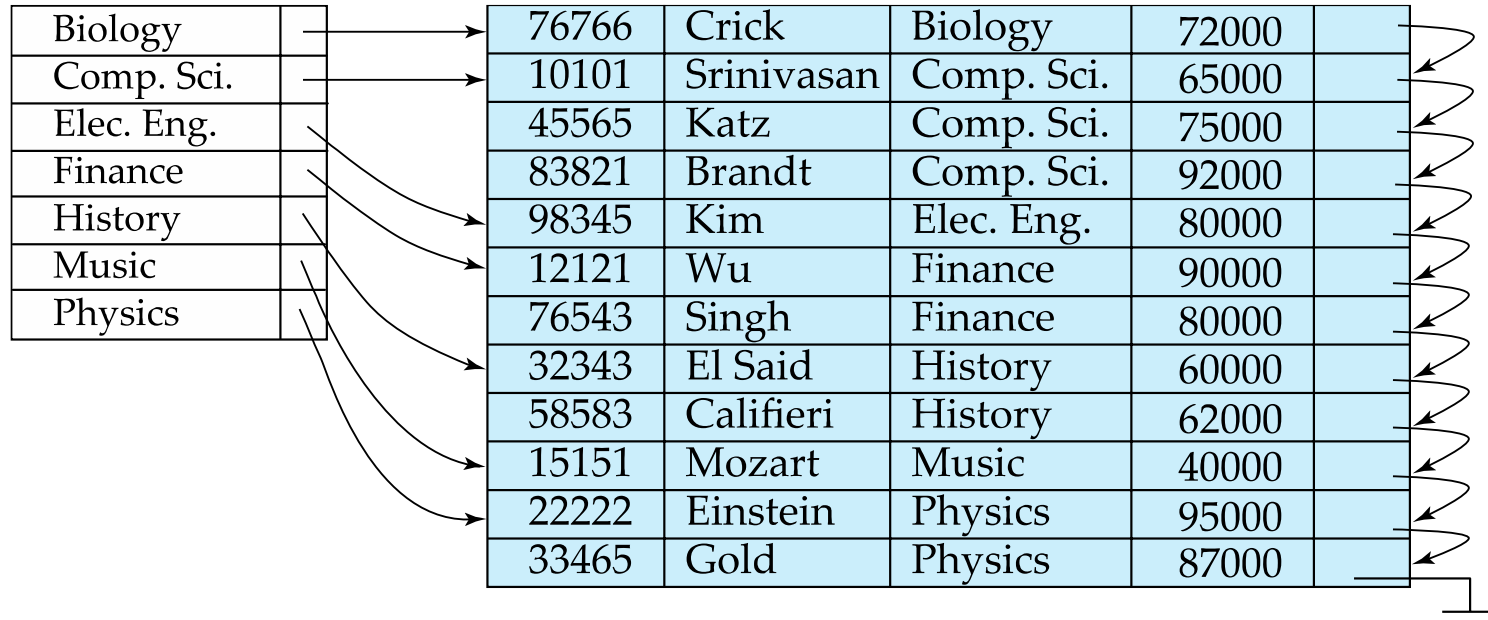
# Dense Index Files

- **Dense index:** Index record appears for every search-key value in the file
  - E.g. index on *ID* attribute of *instructor* relation

10101	→	10101	Srinivasan	Comp. Sci.	65000	↙
12121	→	12121	Wu	Finance	90000	↘
15151	→	15151	Mozart	Music	40000	↙
22222	→	22222	Einstein	Physics	95000	↘
32343	→	32343	El Said	History	60000	↙
33456	→	33456	Gold	Physics	87000	↘
45565	→	45565	Katz	Comp. Sci.	75000	↙
58583	→	58583	Califieri	History	62000	↘
76543	→	76543	Singh	Finance	80000	↙
76766	→	76766	Crick	Biology	72000	↘
83821	→	83821	Brandt	Comp. Sci.	92000	↙
98345	→	98345	Kim	Elec. Eng.	80000	↘

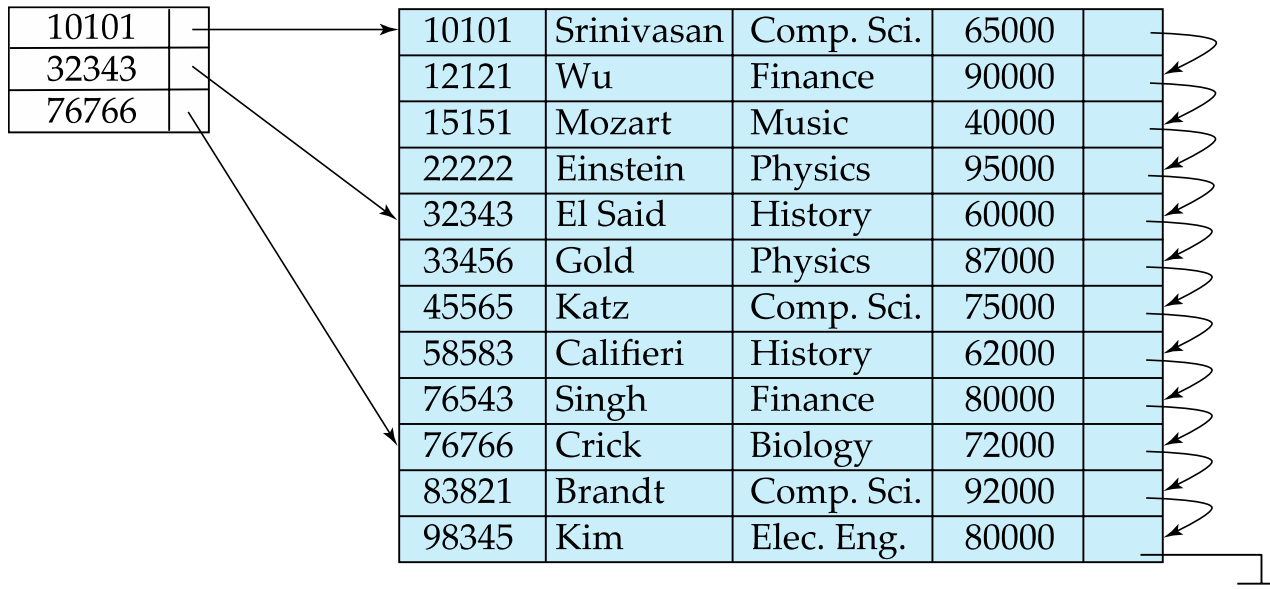
# Dense Index Files

- Dense index on *dept\_name*, with *instructor* file sorted on *dept\_name*



# Sparse Index Files


- **Sparse Index:** Contains index records for only some search-key values
  - Applicable when records are sequentially ordered on search-key
- To locate a record with search-key value  $K$  we:
  - Find index record with largest search-key value  $< K$
  - Search file sequentially starting at the record to which the index record points





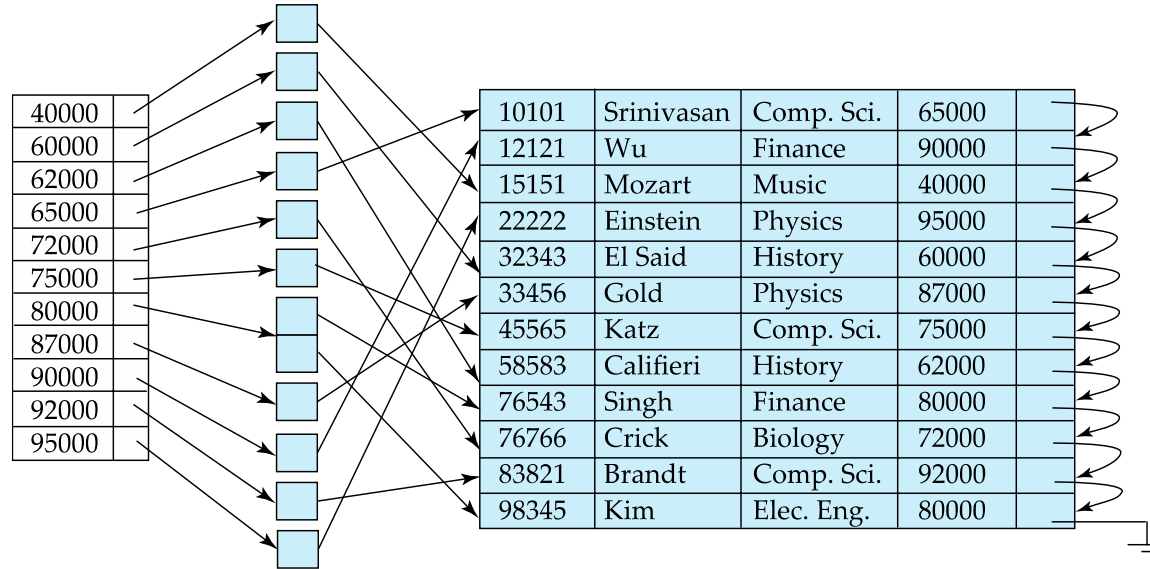
# Sparse Index Files

- Compared to dense indices:
  - Less space and less maintenance overhead for insertions and deletions
  - Generally slower than dense index for locating records
- **Good tradeoff:** Sparse index with an index entry for every block in file, corresponding to least search-key value in the block
- For unclustered index: Sparse index on top of dense index (multilevel index)



# Secondary Indices Example

- Secondary index on *salary* field of instructor



- Index record points to a bucket that contains pointers to all the actual records with that particular search-key value
- Secondary indices have to be dense


# Primary and Secondary Indices

- Indices offer substantial benefits when searching for records
- BUT: Updating indices imposes overhead on database modification, when a file is modified, every index on the file must be updated
- Sequential scan using primary index is efficient, but a sequential scan using a secondary index is expensive
  - Each record access may fetch a new block from disk
  - Block fetch requires about 5 to 10 milliseconds, versus about 100 nanoseconds for memory access

# Multilevel Index

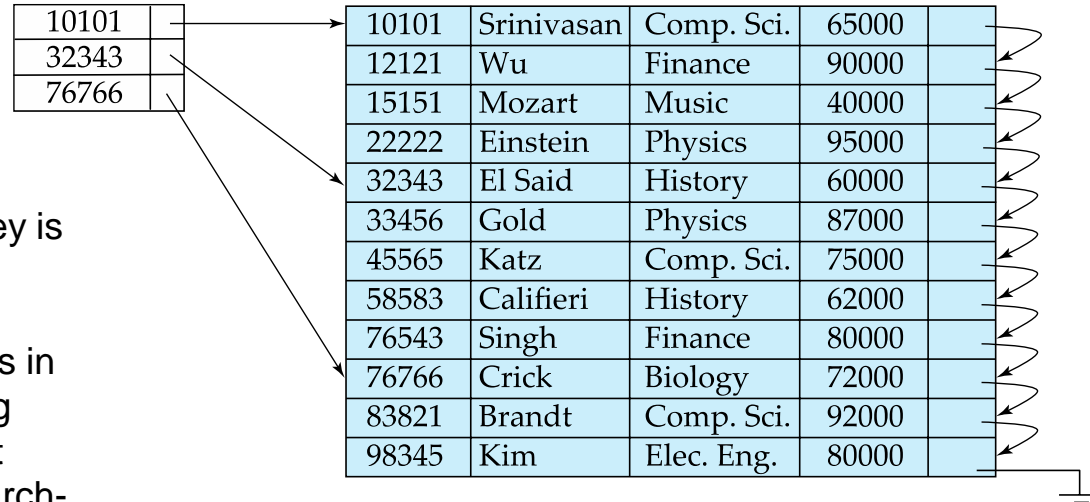
- If primary index does not fit in memory, access becomes expensive
- Solution: Treat primary index kept on disk as a sequential file and construct a sparse index on it
  - **Outer index:** A sparse index of primary index
  - **Inner index:** The primary index file
- If even outer index is too large to fit in main memory, yet another level of index can be created, and so on
- Indices at all levels must be updated on insertion or deletion from the file

# Multilevel Index



# Index Update: Deletion

- If deleted record was the only record in the file with its particular search-key value, the search-key is deleted from the index also



- Single-level index entry deletion:**
  - Dense indices:** Deletion of search-key is similar to file record deletion
  - Sparse indices**
    - If an entry for the search key exists in the index, it is deleted by replacing the entry in the index with the next search-key value in the file (in search-key order)
    - If the next search-key value already has an index entry, the entry is deleted instead of being replaced

# Index Update: Insertion

- **Single-level index insertion:**
  - Perform a lookup using the search-key value of the record to be inserted
  - **Dense indices:** If the search-key value does not appear in the index, insert it
    - Indices are maintained as sequential files
    - Need to create space for new entry, overflow blocks may be required
  - **Sparse indices:** If index stores an entry for each block of the file, no change needs to be made to the index unless a new block is created
    - If a new block is created, the first search-key value appearing in the new block is inserted into the index
- **Multilevel insertion and deletion:** Algorithms are simple extensions of the single-level algorithms

# Secondary Indices

- Frequently, one wants to find all the records whose values in a certain field (which is not the search-key of the primary index) satisfy some condition
  - Example 1: In the *instructor* relation stored sequentially by ID, we may want to find all instructors in a particular department
  - Example 2: As discussed before, but where we want to find all instructors with a specified salary or with salary in a specified range of values
- We can have a secondary index with an index record for each search-key value



# Next Lecture

## 2-3-4 Trees

# Thank you for your attention...

Any question?

**Contact:**

Department of Information Technology, NITK Surathkal, India  
6<sup>th</sup> Floor, Room: 13

**Phone:** +91-9477678768

**E-mail:** [shrutilipi@nitk.edu.in](mailto:shrutilipi@nitk.edu.in)