

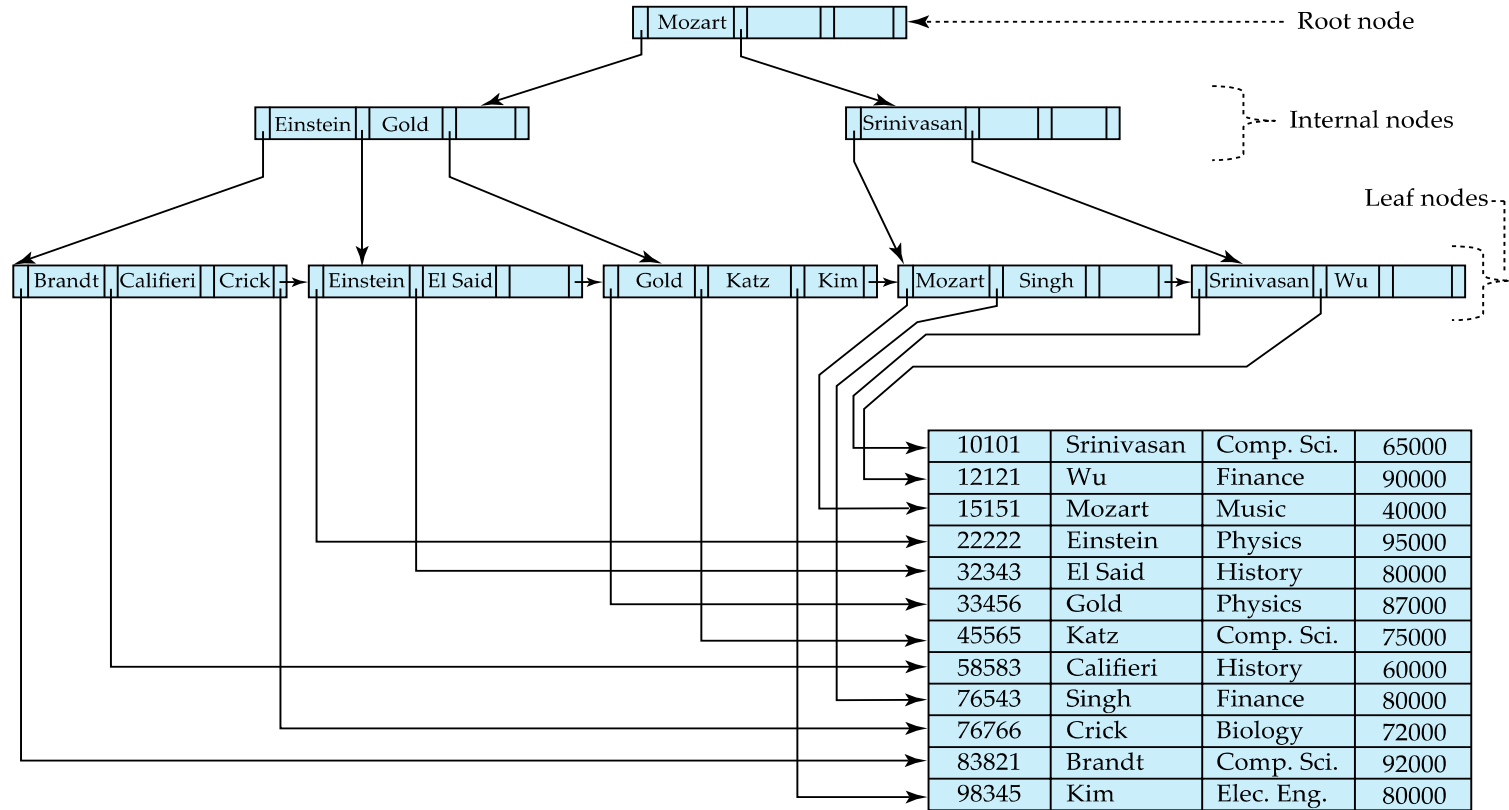


## **B<sup>+</sup> Tree and B Tree**

# B<sup>+</sup> Tree Index Files

- B<sup>+</sup> tree indices are an alternative to indexed-sequential files
- Disadvantage of indexed-sequential files
  - Performance degrades as file grows, since many overflow blocks get created
  - Periodic reorganization of entire file is required
- Advantage of B<sup>+</sup> tree index files:
  - Automatically reorganizes itself with small, local, changes, in the face of insertions and deletions
  - Reorganization of entire file is not required to maintain performance
- (Minor) disadvantage of B<sup>+</sup> trees:
  - Extra insertion and deletion overhead, space overhead
- Advantages of B<sup>+</sup> trees outweigh disadvantages
  - B<sup>+</sup> trees are used extensively

# Example of B+ Tree




# B<sup>+</sup> Tree Index Files

- A B<sup>+</sup> tree is a rooted tree satisfying the following properties:
  - All paths from root to leaf are of the same length
  - Each node that is not a root or a leaf has between  $\lceil n/2 \rceil$  and  $n$  children
  - A leaf node has between  $\lceil (n-1)/2 \rceil$  and  $n-1$  key values
  - Special cases:
    - If the root is not a leaf, it has at least 2 children and maximum  $n$  children
    - If the root is a leaf (that is, there are no other nodes in the tree), it can have between 0 and  $(n-1)$  key values

# B<sup>+</sup> Tree Node Structure

- Typical node



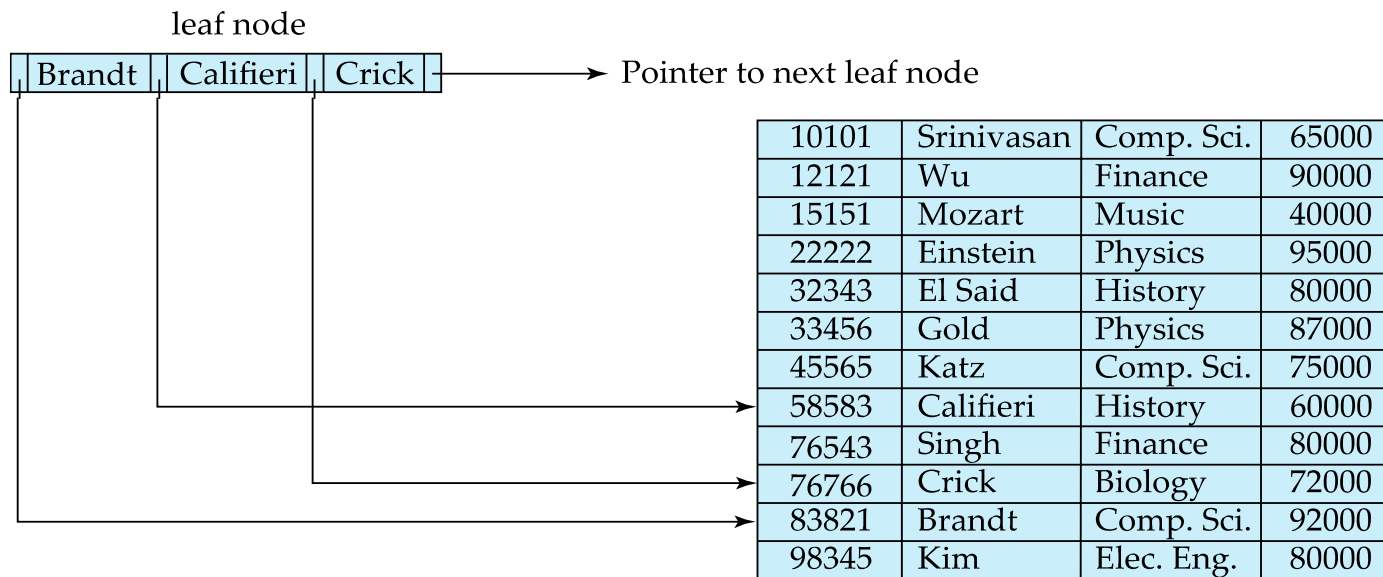
- $K_i$  are the search-key values
  - $P_i$  are pointers to children (for non-leaf nodes) or pointers to records or buckets of records (for leaf nodes)
- The search-keys in a node are ordered

$$K_1 < K_2 < K_3 < \dots < K_{n-1}$$

(Initially assume no duplicate keys, address duplicates later)

# Leaf Nodes in B<sup>+</sup> Trees

- Properties of a leaf node:
  - For  $i = 1, 2, \dots, n-1$ , pointer  $P_i$  points to a file record with search-key value  $K_i$
  - If  $L_i, L_j$  are leaf nodes and  $i < j$ ,  $L_i$ 's search-key values are less than or equal to  $L_j$ 's search-key values
  - $P_n$  points to next leaf node in search-key order




# Non-Leaf Nodes in B<sup>+</sup> Trees

- Non leaf nodes form a multi-level sparse index on the leaf nodes
- For a non-leaf node with  $m$  pointers:
  - All the search-keys in the subtree to which  $P_1$  points are less than  $K_1$
  - For  $2 \leq i \leq n-1$ , all the search-keys in the subtree to which  $P_i$  points have values greater than or equal to  $K_{i-1}$  and less than  $K_i$
  - All the search-keys in the subtree to which  $P_n$  points have values greater than or equal to  $K_{n-1}$
- General structure

$P_1$	$K_1$	$P_2$	$\dots$	$P_{n-1}$	$K_{n-1}$	$P_n$
-------	-------	-------	---------	-----------	-----------	-------

# Example of B<sup>+</sup> Trees

- B<sup>+</sup> tree for *instructor* file ( $n = 6$ )



- Leaf nodes must have between 3 and 5 key values ( $\lceil (n-1)/2 \rceil$  and  $n-1$ , with  $n = 6$ )
- Non-leaf nodes other than root must have between 3 and 6 children ( $\lceil n/2 \rceil$  and  $n$  with  $n = 6$ )
- Root must have at least 2 children



# Observations about B<sup>+</sup> Trees

- Since the inter-node connections are done by pointers, “logically” close blocks need not be “physically” close
- The non-leaf levels of the B<sup>+</sup> tree form a hierarchy of sparse indices
- The B<sup>+</sup> tree contains a relatively small number of levels
  - Level below root has at least  $2 * \lceil n/2 \rceil$  values
  - Next level has at least  $2 * \lceil n/2 \rceil * \lceil n/2 \rceil$  values
  - .. etc.
  - If there are  $K$  search-key values in the file, the tree height is no more than  $\lceil \log_{\lceil n/2 \rceil}(K) \rceil$
  - Thus searches can be conducted efficiently
- Insertions and deletions to the main file can be handled efficiently, as the index can be restructured in logarithmic time

# Queries on B<sup>+</sup> Trees

**function** *find*(*V*)

*C* = *root*

**while** (*C* is not a leaf node) {

    Let *i* be least number s.t.  $V \leq K_i$

**If** no such exists, set *C* = *last non-null pointer in C*

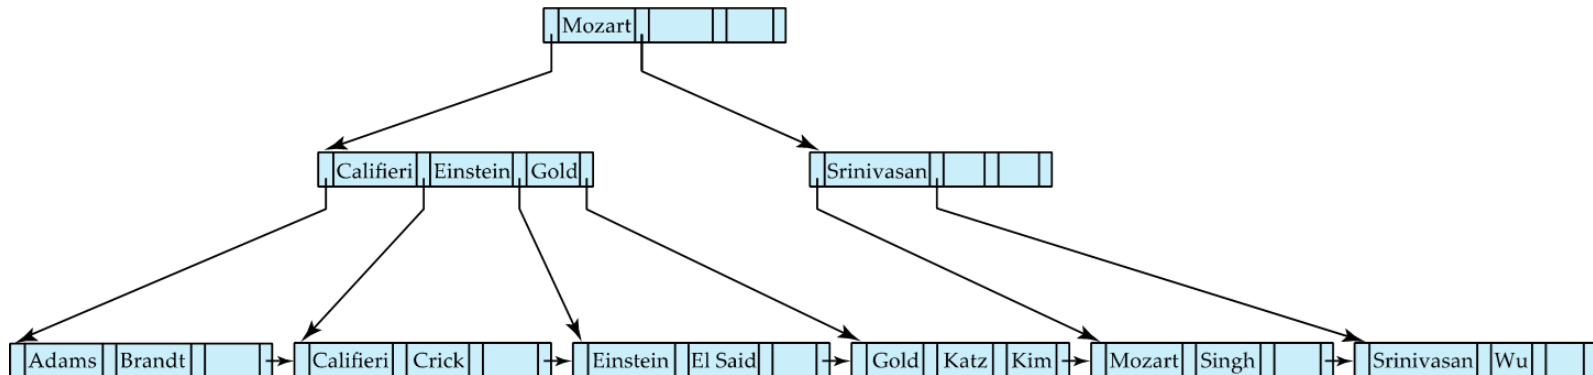
**Else** { **if**  $V = C.K_i$  ) Set *C* =  $P_{i+1}$  **else set** *C* = *C.P<sub>i</sub>* }

}

Let *i* be least value s.t.,  $K_i = V$

**If** there is such a value *i*, follow pointer  $P_i$  to the desired record

**else** no record with search-key value *k* exists



# Handling Duplicates

- With duplicate search keys
  - In both leaf and internal nodes
    - We cannot guarantee that  $K_1 < K_2 < K_3 < \dots < K_{N-1}$
    - But can guarantee  $K_1 \leq K_2 \leq K_3 \leq \dots \leq K_{N-1}$
  - Search-keys in the subtree to which  $P_i$  points
    - Are  $\leq K_i$ , but not necessarily  $< K_i$
    - To see why, suppose same search key value  $V$  is present in two leaf node  $L_i$  and  $L_{i+1}$
    - Then in parent node  $K_i$  must be equal to  $V$

# Handling Duplicates

- We modify find procedure as follows
  - Traverse  $P_i$  even if  $V = K_i$
  - As soon as we reach a leaf node  $C$  check if  $C$  has only search key values less than  $V$
  - If so set  $C =$  right sibling of  $C$  before checking whether  $C$  contains  $V$
- Procedure printAll
  - Uses modified find procedure to find first occurrence of  $V$
  - Traverse through consecutive leaves to find all occurrences of  $V$

# Queries on B<sup>+</sup> Trees


- If there are  $K$  search-key values in the file, the height of the tree is no more than  $\lceil \log_{\lceil N/2 \rceil}(K) \rceil$
- A node is generally the same size as a disk block, typically 4 kilobytes
  - And  $N$  is typically around 100 (40 bytes per index entry)
- With 1 million search key values and  $N = 100$ 
  - At most  $\log_{50}(1,000,000) = 4$  nodes are accessed in a lookup
- Contrast this with a balanced binary tree with 1 million search key values, around 20 nodes are accessed in a lookup
  - Above difference is significant since every node access may need a disk I/O, costing around 20 milliseconds

# Updates on B<sup>+</sup> Trees: Insertion

- Find the leaf node in which the search-key value would appear
- If the search-key value is already present in the leaf node
  - Add record to the file
  - If necessary add a pointer to the bucket
- If the search-key value is not present, then
  - Add the record to the main file (and create a bucket if necessary)
  - If there is room in the leaf node, insert (key-value, pointer) pair in the leaf node
  - Otherwise, split the node (along with the new (key-value, pointer) entry) as discussed in the next slide

# Updates on B<sup>+</sup> Trees: Insertion

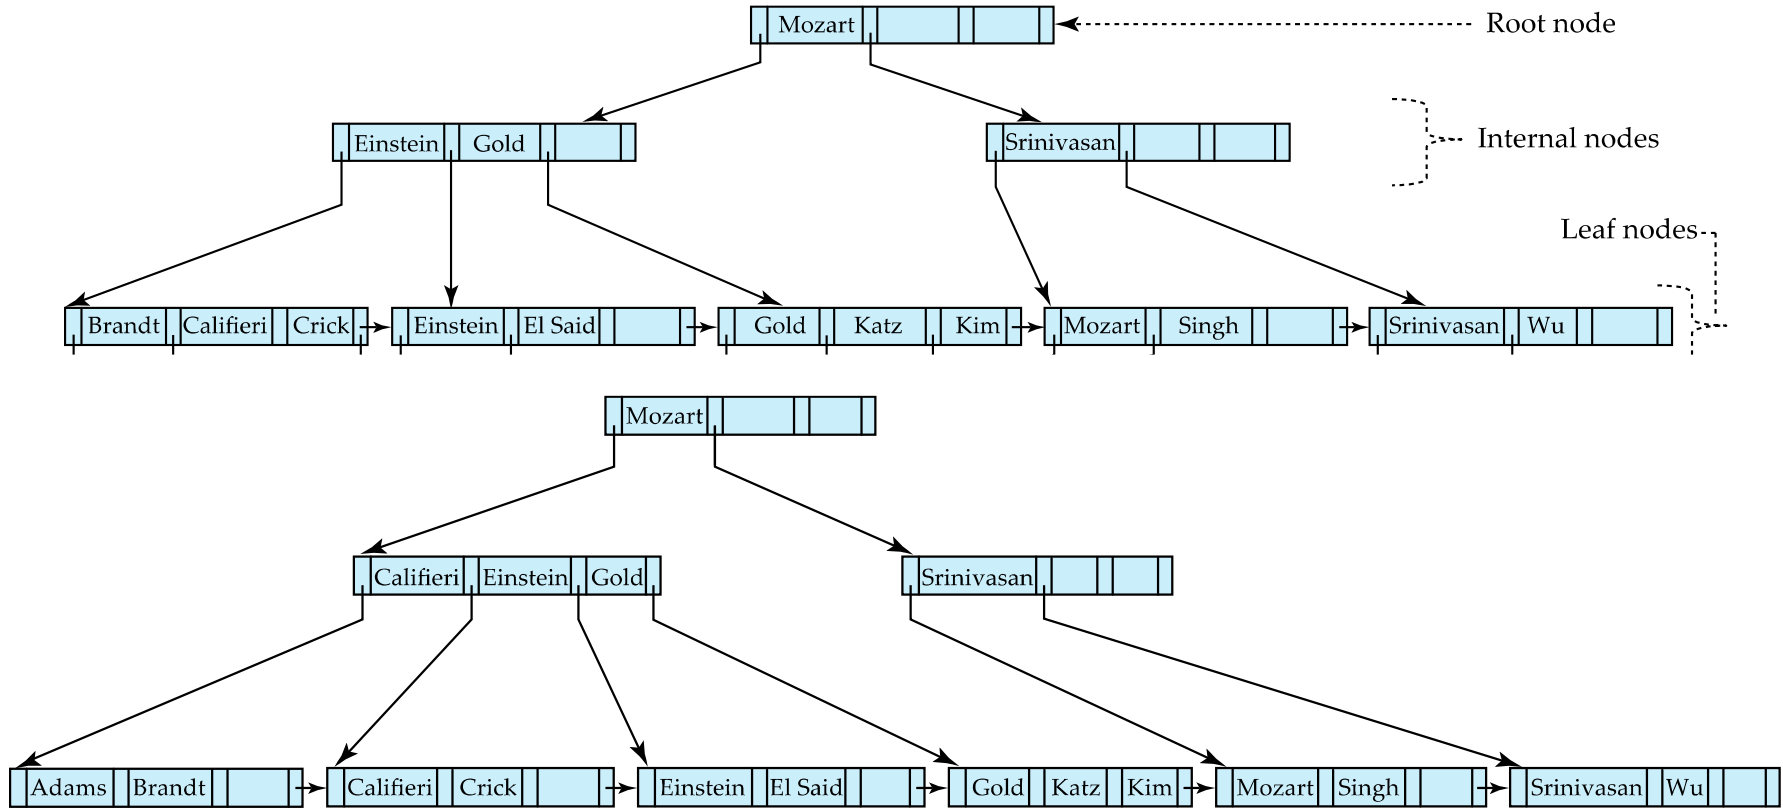
- Splitting a leaf node:
  - Take the  $n$  (search-key value, pointer) pairs (including the one being inserted) in sorted order
  - Place the first  $\lceil n/2 \rceil$  in the original node, and the rest in a new node
  - Let the new node be  $p$ , and let  $k$  be the least key value in  $p$
  - Insert  $(k, p)$  in the parent of the node being split
  - If the parent is full, split it and **propagate** the split further up
- Splitting of nodes proceeds upwards till a node that is not full is found
  - In the worst case the root node may be split increasing the height of the tree by 1



- Result of splitting node containing *Brandt*, *Califieri* and *Crick* on inserting *Adams*
- Next step: insert entry with (*Califieri*, pointer-to-new-node) into parent

# B+ Trees Insertion

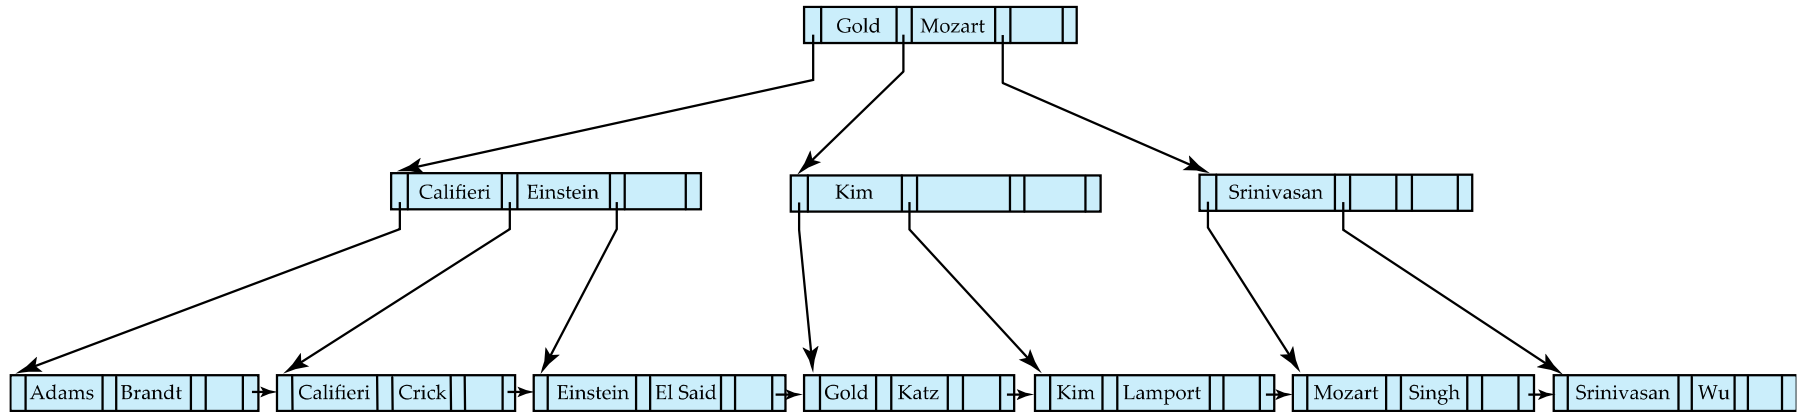
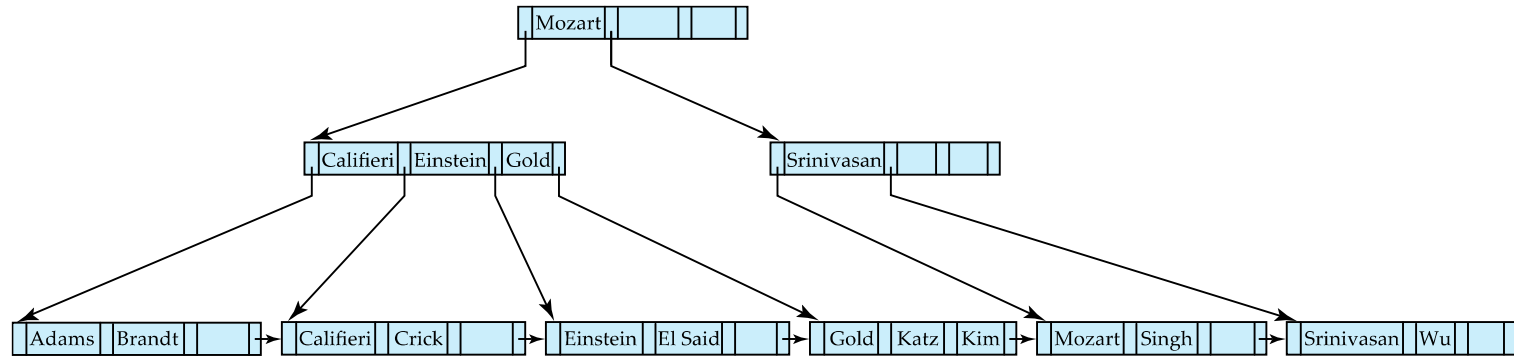
- B+ tree before and after insertion of "Adams"






# B+ Trees Insertion

- B+ tree before and after insertion of "Lamport"



# Insertion in B<sup>+</sup> Trees

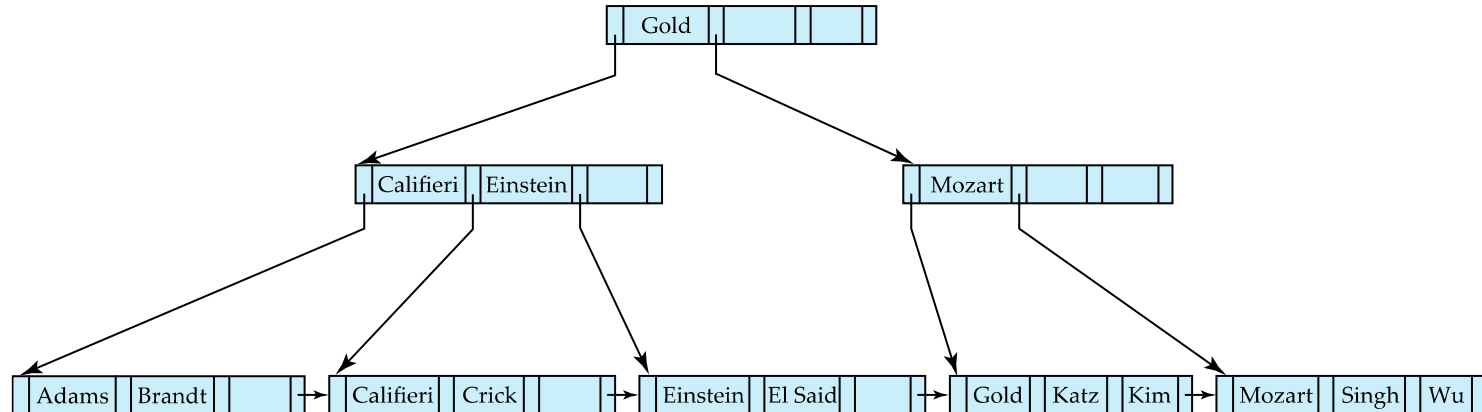
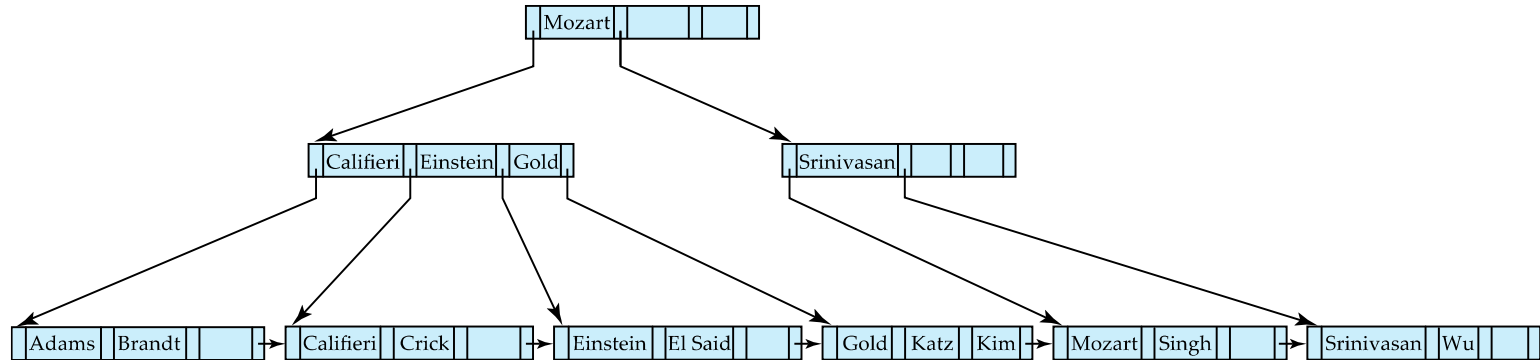
- Splitting a non-leaf node: When inserting  $(k, p)$  into an already full internal node  $N$ 
  - Copy  $N$  to an in-memory area  $M$  with space for  $n + 1$  pointers and  $n$  keys
  - Insert  $(k, p)$  into  $M$
  - Copy  $P_1, K_1, \dots, K_{\lceil n/2 \rceil - 1}, P_{\lceil n/2 \rceil}$  from  $M$  back into node  $N$
  - Copy  $P_{\lceil n/2 \rceil + 1}, K_{\lceil n/2 \rceil + 1}, \dots, K_n, P_{n+1}$  from  $M$  into newly allocated node  $N'$
  - Insert  $(K_{\lceil n/2 \rceil}, N')$  into parent  $N$



**Read pseudocode in book!**

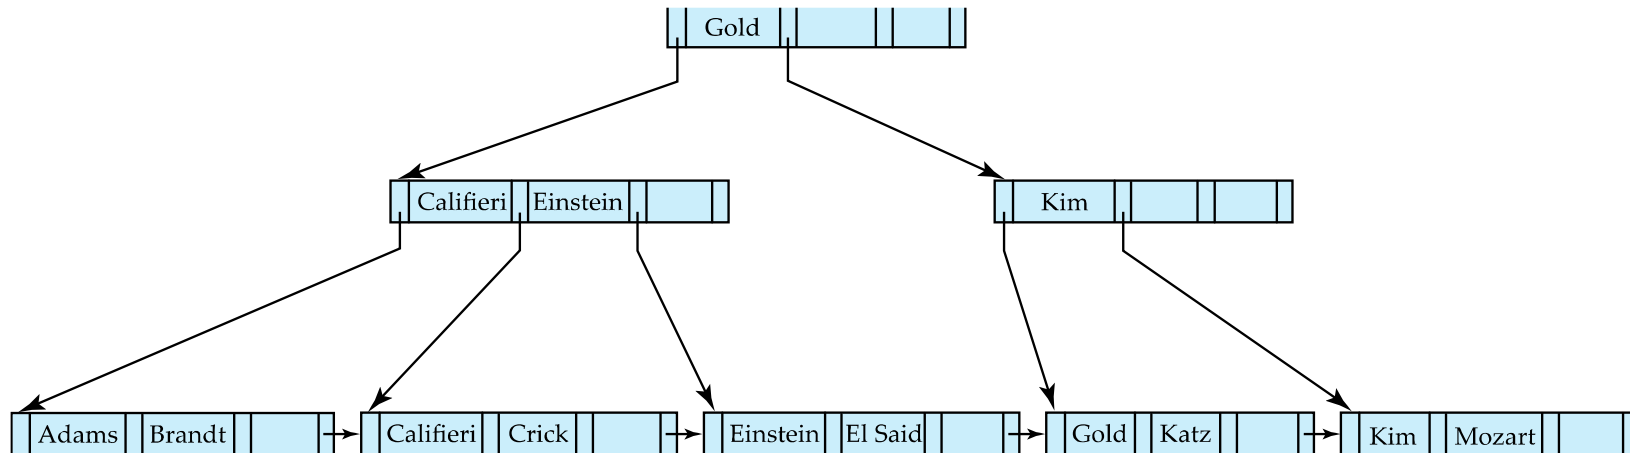
# Examples of B<sup>+</sup> Trees Deletion

- Deleting “Srinivasan” causes merging of under-full leaves



# Examples of B<sup>+</sup> Trees Deletion

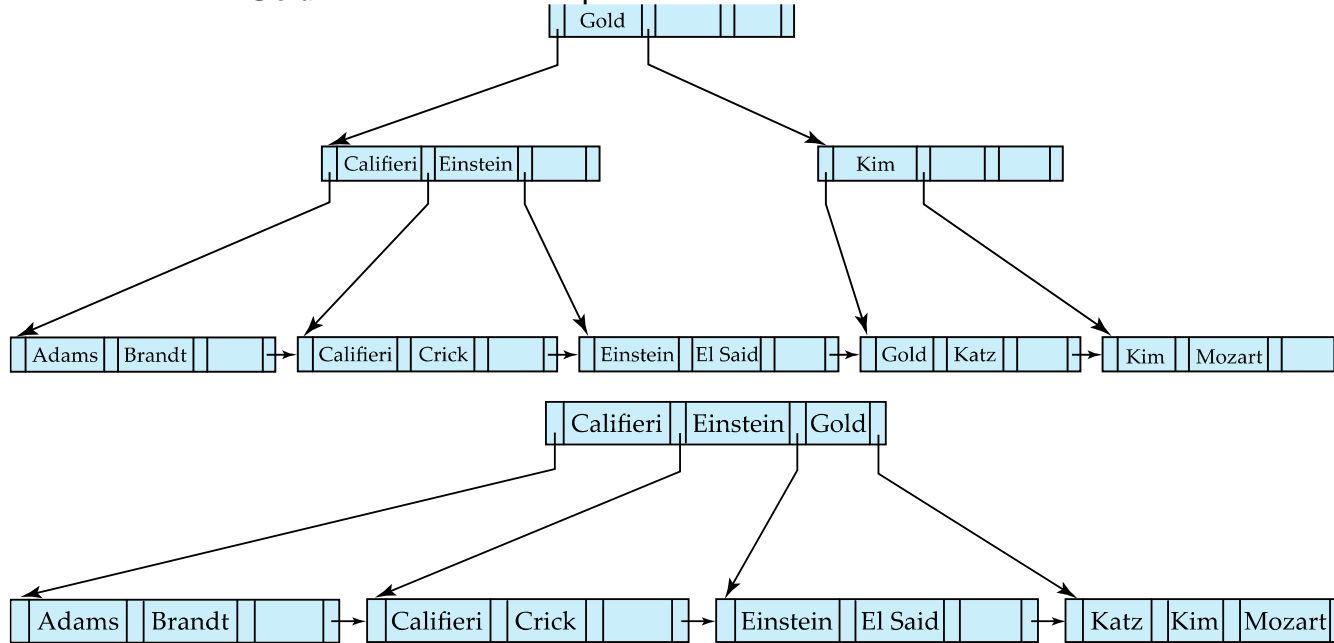
- Deletion of “*Singh*” and “*Wu*” from result of previous example



- Leaf containing *Singh* and *Wu* became underfull, and borrowed a value *Kim* from its left sibling
- Search-key value in the parent changes as a result

# Examples of B<sup>+</sup> Trees Deletion

- Before and after deletion of “*Gold*” from earlier example



- Node with *Gold* and *Katz* became underfull, and was merged with its sibling
- Parent node becomes underfull, and is merged with its sibling
  - Value separating two nodes (at the parent) is pulled down when merging
- Root node then has only one child, and is delete

# Updates on B<sup>+</sup> Trees: Deletion

- Find the record to be deleted, and remove it from the main file and from the bucket (if present)
- Remove (search-key value, pointer) from the leaf node if there is no bucket or if the bucket has become empty
- If the node has too few entries due to the removal, and the entries in the node and a sibling fit into a single node, then ***merge siblings***:
  - Insert all the search-key values in the two nodes into a single node (the one on the left), and delete the other node
  - Delete the pair  $(K_{i-1}, P_i)$ , where  $P_i$  is the pointer to the deleted node, from its parent, recursively using the above procedure

# Updates on B<sup>+</sup> Trees: Deletion

- Otherwise, if the node has too few entries due to the removal, but the entries in the node and a sibling do not fit into a single node, then **redistribute pointers**:
  - Redistribute the pointers between the node and a sibling such that both have more than the minimum number of entries
  - Update the corresponding search-key value in the parent of the node
- The node deletions may cascade upwards till a node which has  $\lceil n/2 \rceil$  or more pointers is found
- If the root node has only one pointer after deletion, it is deleted and the sole child becomes the root


# B<sup>+</sup> Tree File Organization

- Index file degradation problem is solved by using B<sup>+</sup> tree indices
- Data file degradation problem is solved by using B<sup>+</sup> tree file organization
- The leaf nodes in a B<sup>+</sup> tree file organization store records, instead of pointers
- Leaf nodes are still required to be half full
  - Since records are larger than pointers, the maximum number of records that can be stored in a leaf node is less than the number of pointers in a non-leaf node
- Insertion and deletion are handled in the same way as insertion and deletion of entries in a B<sup>+</sup> tree index



# B<sup>+</sup> Tree File Organization

- Example of B<sup>+</sup> tree file organization



- Good space utilization important since records use more space than pointers
- To improve space utilization, involve more sibling nodes in redistribution during splits and merges
  - Involving 2 siblings in redistribution (to avoid split / merge where possible) results in each node having at least entries

# Other Issues in Indexing

## Record relocation and secondary indices


- If a record moves, all secondary indices that store record pointers have to be updated
- Node splits in B<sup>+</sup> tree file organizations become very expensive
- *Solution*: Use search key of B<sup>+</sup> tree file organization instead of record pointer in secondary index
  - Extra traversal of primary index to locate record
    - Higher cost for queries, but node splits are cheap
  - Add record-id if primary-index search key is non-unique

# Indexing Strings


- Variable length strings as keys
  - Variable fan-out
  - Use space utilization as criterion for splitting, not number of pointers
- **Prefix compression**
  - Key values at internal nodes can be prefixes of full key
    - Keep enough characters to distinguish entries in the subtrees separated by the key value
      - E.g., “Silas” and “Silberschatz” can be separated by “Silb”
  - Keys in leaf node can be compressed by sharing common prefixes

# B-Tree Index Files

- Similar to B<sup>+</sup> tree, but B-tree allows search-key values to appear only once; eliminates redundant storage of search keys
- Search keys in non-leaf nodes appear nowhere else in the B-tree; an additional pointer field for each search key in a non-leaf node must be included
- Generalized B-tree leaf node



(a)

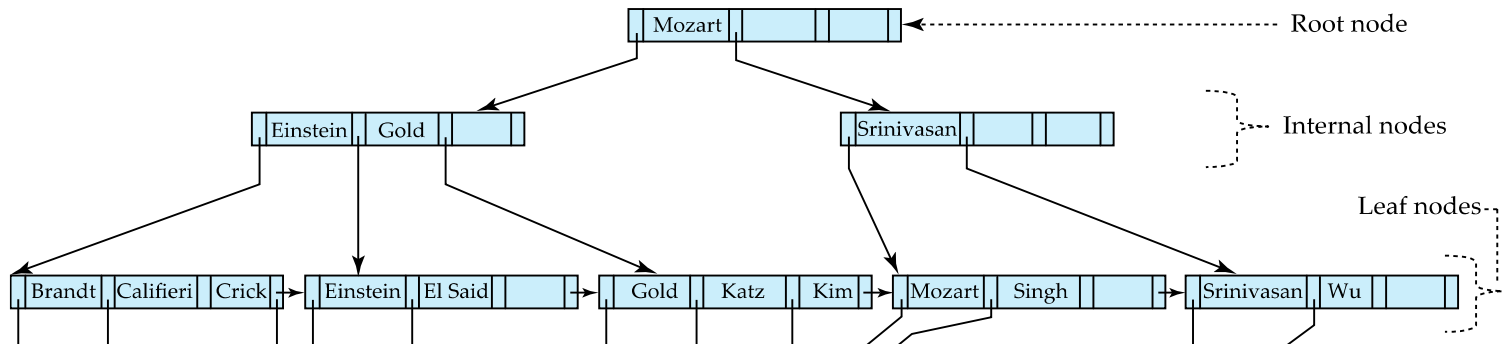
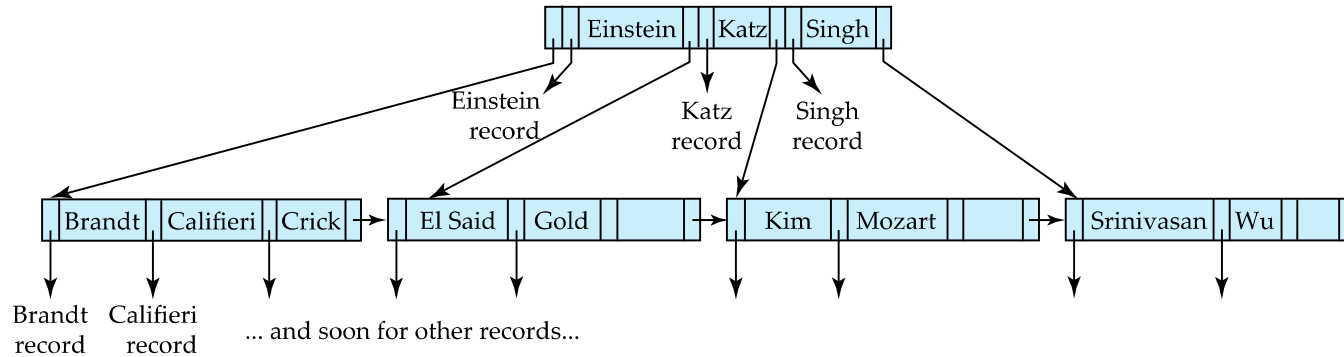


(b)

- Non-leaf node pointers  $B_i$  are the bucket or file record pointers

# B-Tree Index File Example

- B-tree (above) and B<sup>+</sup> tree (below) on same data



# B-Tree Index File Example

- Advantages of B-tree indices:
  - May use less tree nodes than a corresponding B<sup>+</sup> tree
  - Sometimes possible to find search-key value before reaching leaf node
- Disadvantages of B-tree indices:
  - Only small fraction of all search-key values are found early
  - Non-leaf nodes are larger, so fan-out is reduced
  - Thus, B-trees typically have greater depth than corresponding B<sup>+</sup> tree
  - Insertion and deletion more complicated than in B<sup>+</sup> tree
  - Implementation is harder than B<sup>+</sup> tree
- Typically, advantages of B-trees do not outweigh disadvantages

# Next Lecture

## Hashing

# Thank you for your attention...

Any question?

**Contact:**

Department of Information Technology, NITK Surathkal, India  
6<sup>th</sup> Floor, Room: 13

**Phone:** +91-9477678768

**E-mail:** [shrutilipi@nitk.edu.in](mailto:shrutilipi@nitk.edu.in)