

Scheduling Algorithms

Shortest Job First and Shortest Remaining Time First

Non-preemptive vs. Preemptive Scheduling

Under *non-preemptive scheduling*, each running process keeps the CPU until it completes or it switches to the waiting (blocked) state

Under *preemptive scheduling*, a running process may be also forced to release the CPU even though it is neither completed nor blocked.

- In time-sharing systems, when the running process reaches the end of its time *quantum*

Algorithm 3: Non-Preemptive Shortest Job First (SJF)


Characteristics:

- Always assign job with shortest next CPU burst to CPU

Example:

<u>Process</u>	<u>Arrival</u>	<u>Burst</u>
P1	0	7
P2	2	4
P3	4	1
P4	5	4

Gantt Chart




Avg wait time: $(0 + 6 + 3 + 7) / 4 = 4$ ms

Example for Non-Preemptive SJF

	<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
	P_1	0.0	7
	P_2	2.0	4
	P_3	4.0	1
	P_4	5.0	4


SJF (non-preemptive)



Example for Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

Once P_1 has completed the queue now holds P_2 , P_3 and P_4



P_3 gets the CPU first since it is the shortest. P_2 then P_4 get the CPU in turn (based on arrival time)

Algorithm 4: Preemptive Shortest Job First (SJF)- Shortest-Remaining-Time-First (SRTF)


Characteristics:

- Like Shortest Job First, but running job can be preempted.

Example:

<u>Process</u>	<u>Arrival</u>	<u>Burst</u>
P1	0	7
P2	2	4
P3	4	1
P4	5	4

Gantt Chart



Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$ ms


Example for Preemptive SJF (SRTF)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

Time 0 – P_1 gets the CPU Ready = $[(P_1, 7)]$

Time 2 – P_2 arrives – CPU has P_1 with time=5, Ready = $[(P_2, 4)]$ – P_2 gets the CPU

Time 4 – P_3 arrives – CPU has P_2 with time = 2, Ready = $[(P_1, 5), (P_3, 1)]$ – P_3 gets the CPU




Example for Preemptive SJF (SRTF)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

Time 5 – P_3 completes and P_4 arrives - Ready = $[(P_1,5),(P_2,2),(P_4,4)]$ – P_2 gets the CPU

Time 7 – P_2 completes – Ready = $[(P_1,5),(P_4,4)]$ – P_4 gets the CPU


Time 11 – P_4 completes, P_1 gets the CPU



Example for Preemptive SJF (SRTF)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

SJF (preemptive)



$$\text{Average waiting time} = (9 + 1 + 0 + 2)/4 = 3$$

Implementing SJF

SJF is great, but how do we implement it?

- We don't know a priori how long a job's burst time is
- We have to try to *predict* the burst time

Priority-Based Scheduling

A priority number (integer) is associated with each process

The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority).

- Preemptive
- Non-preemptive

SJF is a priority scheme with the priority the remaining time.

Example for Non Preemptive Priority-based Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

