

An Introduction to Device Drivers

Sarah Diesburg
COP 5641 / CIS 4930

Introduction

■ Device drivers

- ❑ Black boxes to hide details of hardware devices
 - ❑ Use standardized calls
 - Independent of the specific driver
 - ❑ Main role
 - Map standard calls to device-specific operations
 - ❑ Can be developed separately from the rest of the kernel
 - Plugged in at runtime when needed
-

The Role of the Device Driver

- Implements the *mechanisms* to access the hardware
 - E.g., show a disk as an array of data blocks
 - Does not force particular *policies* on the user
 - Examples
 - Who many access the drive
 - Whether the drive is accessed via a file system
 - Whether users may mount file systems on the drive
-

Policy-Free Drivers

- A common practice
 - Support for synchronous/asynchronous operation
 - Be opened multiple times
 - Exploit the full capabilities of the hardware
 - Easier user model
 - Easier to write and maintain
 - To assist users with policies, release device drivers with user programs
-

Splitting the Kernel

- Process management
 - Creates, destroys processes
 - Supports communication among processes
 - Signals, pipes, etc.
 - Schedules how processes share the CPU
 - Memory management
 - Virtual addressing
-

Splitting the Kernel

■ File systems

- Everything in UNIX can be treated as a file
- Linux supports multiple file systems

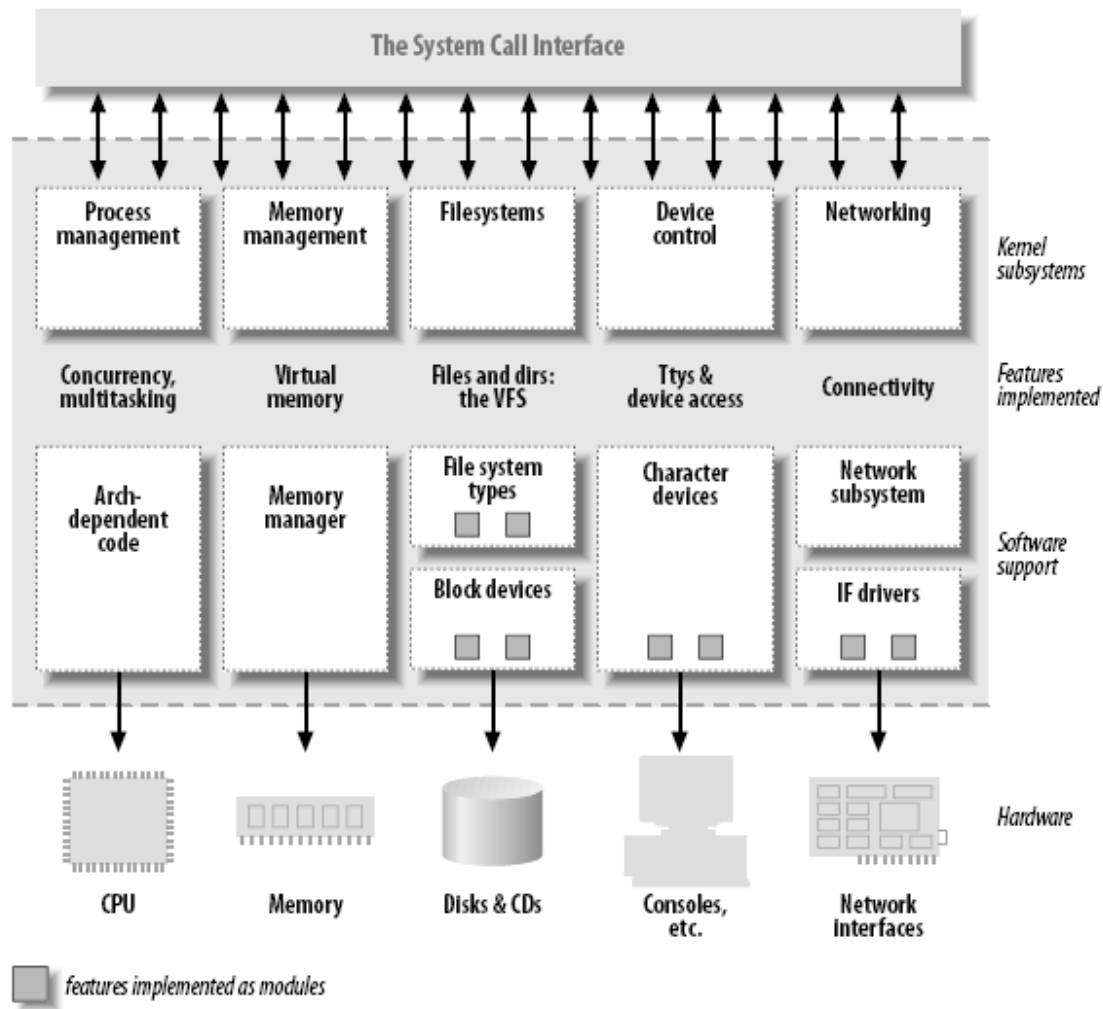
■ Device control

- Every system operation maps to a physical device
 - Few exceptions: CPU, memory, etc.

■ Networking

- Handles packets
- Handles routing and network address resolution issues

Splitting the Kernel



Loadable Modules

- The ability to add and remove kernel features at runtime
 - Each unit of extension is called a *module*
 - Use `insmod` program to add a kernel module
 - Use `rmmmod` program to remove a kernel module
-

Classes of Devices and Modules

- Character devices
 - Block devices
 - Network devices
 - Others
-

Character Devices

- Abstraction: a stream of bytes
 - Examples
 - Text console (`/dev/console`)
 - Serial ports (`/dev/ttyS0`)
 - Usually supports **open**, **close**, **read**, **write**
 - Accessed sequentially (in most cases)
 - Might not support file seeks
 - Exception: frame grabbers
 - Can access acquired image using **mmap** or **lseek**

Block Devices

- Abstraction: array of storage blocks
 - However, applications can access a block device in bytes
 - Block and char devices differ only at the kernel level
 - A block device can host a file system
-

Network Devices

- Abstraction: data packets
 - Send and receive packets
 - Do not know about individual connections
 - Have unique names (e.g., **eth0**)
 - Not in the file system
 - Support protocols and streams related to packet transmission (i.e., no **read** and **write**)
-

Other Classes of Devices

- Examples that do not fit to previous categories:
 - ❑ USB
 - ❑ SCSI
 - ❑ FireWire
 - ❑ MTD
-

File System Modules

- Software drivers, not device drivers
 - Serve as a layer between user API and block devices
 - Intended to be device-independent
-

Security Issues

- Deliberate vs. incidental damage
 - Kernel modules present possibilities for both
 - System does only rudimentary checks at module load time
 - Relies on limiting privilege to load modules
 - And trusts the driver writers
 - Driver writer must be on guard for security problems
-

Security Issues

- Do not define security policies
 - Provide mechanisms to enforce policies
 - Be aware of operations that affect global resources
 - Setting up an interrupt line
 - Could damage hardware
 - Setting up a default block size
 - Could affect other users
-

Security Issues

- Beware of bugs
 - Buffer overrun
 - Overwriting unrelated data
 - Treat input/parameters with utmost suspicion
 - Uninitialized memory
 - Kernel memory should be zeroed before being made available to a user
 - Otherwise, information leakage could result
 - Passwords

Security Issues

- Avoid running kernels compiled by an untrusted friend
 - Modified kernel could allow anyone to load a module

Version Numbering

- Every software package used in Linux has a release number
 - You need a particular version of one package to run a particular version of another package
 - Prepackaged distribution contains matching versions of various packages
-

Version Numbering

- Different throughout the years
 - After version 1.0 but before 3.0
 - <major>.<minor>.<release>.<bugfix>
 - Time based releases (after two to three months)
 - 3.x
 - Moved to 3.0 to commemorate 20th anniversary of Linux
 - <version>.<release>.<bugfix>
 - <https://lkml.org/lkml/2011/5/29/204>
-

License Terms

- GNU General Public License (GPL2)
 - GPL allows anybody to redistribute and sell a product covered by GPL
 - As long as the recipient has access to the source
 - And is able to exercise the same rights
 - Any software product derived from a product covered by the GPL be released under GPL
-

License Terms

- If you want your code to go into the mainline kernel
 - Must use a GPL-compatible license

Joining the Kernel Development Community

- The central gathering point
 - Linux-kernel mailing list
 - <http://www.tux.org/lkml>
- Chapter 20 of LKM further discusses the community and accepted coding style